

# Assignment 2 - Deadline: Oct 9, 2024, Wed 11pm

## DSAI 510 Fall 2024

Complete the assignment below and upload **both the .ipynb file and its pdf** to <https://moodle.boun.edu.tr> (<https://moodle.boun.edu.tr>) by the deadline given above. The submission page on Moodle will close automatically after this date and time.

To make a pdf, this may work: Hit CMD+P or CTRL+P, and save it as PDF. You may also use other options from the File menu.

```
In [1]: # Run this cell first

import pandas as pd
import numpy as np

# Set the display option to show all rows scrolling with a slider
pd.set_option('display.max_rows', None)
# To disable this, run the line below:
# pd.reset_option('display.max_rows')
```

## Note:

In the problems below, if it asks, "show the number of records that are nonzero", the answer is a number; so you don't need to show the records themselves. But if it asks, "show the records with NaN", it wants you to print those records (rows) containing NAN and other entries, not asking how many such records there are. So be careful about what you're asked.

## Problem 1 (10 pts)

- Load **Electric\_Vehicle\_Population\_Data-modified1.csv** and **Electric\_Vehicle\_Population\_Data-modified2.csv** into pandas dataframes as df1 and df2.
- Inspect the first and last five records with `head()` and `tail()` for both dataframes.
- Use `len()` and `print()` [or `display()`] to show how many records each dataframe contains.
- Use `info()` to get a summary of both dataframes.
- Combine df1 and df2 into a new dataframe called df3 by using `concat()` and print the number of records in the new dataframe df3.
- Find and print the number of duplicate records by using `duplicated()` and `sum()`.
- Drop duplicates, save the new dataframe as dfALL and then print the number of records in dfALL.

```
In [3]: # Break your computations into multiple cells.

df1 = pd.read_csv("Electric_Vehicle_Population_Data-modified1.csv")
df2 = pd.read_csv("Electric_Vehicle_Population_Data-modified2.csv")

display(df1.head())
display(df1.tail())
display(df2.head())
display(df1.tail())
```

	VIN (1-10)	County	City	State	Postal Code	Model Year	Make	Model	Electric Vehicle Type
0	KM8K33AGXL	King	Seattle	WA	98103	2020	HYUNDAI	KONA	Battery Electric Vehicle (BEV)
1	1C4RJYB61N	King	Bothell	WA	98011	2022	JEEP	GRAND CHEROKEE	Plug-in Hybrid Electric Vehicle (PHEV)
2	1C4RJYD61P	Yakima	Yakima	WA	98908	2023	JEEP	GRAND CHEROKEE	Plug-in Hybrid Electric Vehicle (PHEV)
3	5YJ3E1EA7J	King	Kirkland	WA	98034	2018	TESLA	MODEL 3	Battery Electric Vehicle (BEV)
4	WBY7Z8C5XJ	Thurston	Olympia	WA	98501	2018	BMW	I3	Plug-in Hybrid Electric Vehicle (PHEV)

	VIN (1-10)	County	City	State	Postal Code	Model Year	Make	Model	Elk Ve
108441	WBY8P8C55K	King	Seattle	WA	98105	2019	BMW	I3	P H El V (F
108442	YV4H60CF3R	Pierce	Graham	WA	98338	2024	VOLVO	XC90	P H El V (F
108443	1FADP5CU7F	Snohomish	Monroe	WA	98272	2015	FORD	C-MAX	P H El V (F
108444	1G1FZ6S07L	Snohomish	Bothell	WA	98012	2020	CHEVROLET	BOLT EV	B El V
108445	5YJ3E1EB1M	Grant	Moses Lake	WA	98837	2021	TESLA	MODEL 3	B El V

	VIN (1-10)	County	City	State	Postal Code	Model Year	Make	Model	Ele Ve
0	1FMCU0EZ1N	Chelan	Wenatchee	WA	98801.0	2022	FORD	ESCAPE	Pl H Ele Ve (P
1	5YJ3E1EB9K	Snohomish	Arlington	WA	98223.0	2019	TESLA	MODEL 3	Ba Ele Ve (
2	5YJSA1E57N	King	Woodinville	WA	98072.0	2022	TESLA	MODEL S	Ba Ele Ve (
3	5YJ3E1EB4J	Snohomish	Snohomish	WA	98290.0	2018	TESLA	MODEL 3	Ba Ele Ve (
4	KL8CK6S00F	Whatcom	Bellingham	WA	98225.0	2015	CHEVROLET	SPARK	Ba Ele Ve (

	VIN (1-10)	County	City	State	Postal Code	Model Year	Make	Model	Elk Ve
108441	WBY8P8C55K	King	Seattle	WA	98105	2019	BMW	I3	P H El V (F
108442	YV4H60CF3R	Pierce	Graham	WA	98338	2024	VOLVO	XC90	P H El V (F
108443	1FADP5CU7F	Snohomish	Monroe	WA	98272	2015	FORD	C-MAX	P H El V (F
108444	1G1FZ6S07L	Snohomish	Bothell	WA	98012	2020	CHEVROLET	BOLT EV	B El V
108445	5YJ3E1EB1M	Grant	Moses Lake	WA	98837	2021	TESLA	MODEL 3	B El V

```
In [ ]: # Comment your code in your own words (not GPT) unless the line is obvious.
```

```
In [5]: print(len(df1))
print(len(df2))
```

```
108446
50484
```

```
In [7]: df1.info()
print()
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 108446 entries, 0 to 108445
Data columns (total 17 columns):
#   Column                                Non-Null C
ount    Dtype
---  -
----  -
```

```

0    VIN (1-10)                                108446 non
-null object
1    County                                    108446 non
-null object
2    City                                      108446 non
-null object
3    State                                    108446 non
-null object
4    Postal Code                              108446 non
-null int64
5    Model Year                               108446 non
-null int64
6    Make                                     108446 non
-null object
7    Model                                    108446 non
-null object
8    Electric Vehicle Type                    108446 non
-null object
9    Clean Alternative Fuel Vehicle (CAFV) Eligibility 108446 non
-null object
10   Electric Range                           108446 non
-null int64
11   Base MSRP                               108446 non
-null int64
12   Legislative District                     108407 non
-null float64
13   DOL Vehicle ID                           108446 non
-null int64
14   Vehicle Location                         108445 non
-null object
15   Electric Utility                         108446 non
-null object
16   2020 Census Tract                       108446 non
-null int64
dtypes: float64(1), int64(6), object(10)
memory usage: 14.1+ MB

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50484 entries, 0 to 50483
Data columns (total 17 columns):
#    Column                                Non-Null C
ount  Dtype
---  ---
----  ---
0    VIN (1-10)                            50484 non-
null object
1    County                                50481 non-
null object
2    City                                  50481 non-
null object
3    State                                50484 non-
null object
4    Postal Code                           50481 non-

```

```

null    float64
   5    Model Year                                50484 non-
null    int64
   6    Make                                      50484 non-
null    object
   7    Model                                    50484 non-
null    object
   8    Electric Vehicle Type                    50484 non-
null    object
   9    Clean Alternative Fuel Vehicle (CAFV) Eligibility 50484 non-
null    object
  10    Electric Range                            50484 non-
null    int64
  11    Base MSRP                                50484 non-
null    int64
  12    Legislative District                      50170 non-
null    float64
  13    DOL Vehicle ID                          50484 non-
null    int64
  14    Vehicle Location                        50478 non-
null    object
  15    Electric Utility                        50481 non-
null    object
  16    2020 Census Tract                       50481 non-
null    float64
dtypes: float64(3), int64(4), object(10)
memory usage: 6.5+ MB

```

```
In [9]: df3 = pd.concat([df1,df2])
```

```
print(len(df3))
```

```
158930
```

```
In [11]: print(df3.duplicated().sum())
```

```
8448
```

```
In [13]: dfALL = df3.drop_duplicates()
```

```
print(len(dfALL))
```

```
150482
```

## Problem 2 (10 pts)

- Make a new dataframe, keep the columns **Model Year**, **Make**, **Model**, **Electric Range**, **Vehicle Location**, and drop all other columns.
- Change the column name **Model Year** into **Year**.
- Show the record with index number 10.
- As you see, the **Vehicle Location** shows the coordinates in the format "POINT (-122.20264 47.6785)". Here the first number (-122.xxx) is the longitude and second number is the latitude. Make two new columns **Latitude** and **Longitude**, carry the numbers to these columns by using `str` method from pandas. Finally change the type of **Latitude** and **Longitude** into float if they're not already. Finally, drop the column **Vehicle Location**.

```
In [15]: # Take wanted columns as below
columns = [
#     'VIN (1-10)', 'County', 'City', 'State', 'Postal Code',
#     'Model Year', 'Make', 'Model', 'Electric Range', 'Vehicle Location',
#     'Electric Vehicle Type', 'Clean Alternative Fuel Vehicle (CAF
#     V) Eligibility',
#     'Base MSRP', 'Legislative District', 'DOL Vehicle ID',
#     'Electric Utility', '2020 Census Tract'
]

# Create new DataFrame from above columns
df = dfALL[columns]
```

```
In [17]: # Use str method to change "Model Year" to "Year"
df.columns = df.columns.str.replace("Model Year", "Year")

df.columns
```

```
Out[17]: Index(['Year', 'Make', 'Model', 'Electric Range', 'Vehicle Location'], dtype='object')
```

```
In [19]: # .iloc[] stands to reach values by index
df.iloc[10]
```

```
Out[19]: Year                2018
Make                TESLA
Model                MODEL 3
Electric Range                215
Vehicle Location    POINT (-122.20264 47.6785)
Name: 10, dtype: object
```



```

In [ ]: # Alternative solution to extract Longitude and Latitude

def extract_points(input_string):

    long = np.nan
    lati = np.nan

    # Control if a POINT string is acceptable format and non-null
    if input_string and " " in str(input_string):
        # Replace the "POINT (" with nothing and ")" with nothing
        cleaned_string = str(input_string).replace("POINT (", "").replace(")", "")

        # split the rest of the string and extract Longitude and Latitude
        cleaned_string = cleaned_string.split(" ")

        long = cleaned_string[0]
        lati = cleaned_string[1]

    # Return coordinates as below format
    return {"Longitude":float(long), "Latitude":float(lati)}

# df["Vehicle Location"].head()

# Change the format of "Vehicle Location" column
df["Vehicle Location"] = df["Vehicle Location"].apply(extract_points)

# Create "Longitude" and "Latitude" columns based on "Vehicle Location" column
for i in df["Vehicle Location"][0].keys():
    df[i] = df["Vehicle Location"].apply(lambda d: d[i])

df = df.drop(columns=['Vehicle Location'])

```

```

In [21]: df = df.copy()

# Create Longitude and Latitude columns and extract the data from
# "Vehicle Location" column

# matches the string "POINT (" followed by two space-separated values,
# each consisting of one or more characters that are not spaces,
# ending with a closing parenthesis.
df['Longitude'] = df['Vehicle Location'].str.extract(r'POINT \(([^\s]+) ([^\s]+)\)')[0].astype(float)
df['Latitude'] = df['Vehicle Location'].str.extract(r'POINT \(([^\s]+) ([^\s]+)\)')[1].astype(float)

df = df.drop(columns=['Vehicle Location'])

```

## Problem 3 (10 pts)

- The file **EV\_prices.csv** contains prices for various makes, models, and years of cars. Load this file into a dataframe. Rename the column **Model Year** to **Year**.
- We want to add a new column **Price** to our dataframe from the previous problem. This column will include the price of the car for the corresponding make, model and year if this information is available in the file **EV\_prices.csv**. If not available, we'll still keep the record but the entry for price will be empty, NA, None or NaN. To achieve this, merge the dataframe from the previous problem with the dataframe containing the data from **EV\_prices.csv**. Think carefully and decide if you need to merge with 'inner' or 'outer' method. At the end, we should have these columns in the merged dataframe: **Year**, **Make**, **Model**, **Electric Range**, **Latitude**, **Longitude** and **Price**. Again, the **Price** column will have numbers only for some records, but it will be empty or NaN for the rest.
- Next, show the number of records which has price information in the **Price** column. Hint: You can use a one-liner containing `len()` and `dropna()` together.

```
In [23]: df_EV = pd.read_csv("EV_prices.csv")

print(df_EV.columns)

df_EV.columns = df_EV.columns.str.replace("Model Year", "Year")

print(df_EV.columns)

Index(['Model Year', 'Make', 'Model', 'Price'], dtype='object')
Index(['Year', 'Make', 'Model', 'Price'], dtype='object')
```

```
In [25]: df_merged = df.merge(df_EV, on=["Year", "Make", "Model"], how="left")

print(df_merged.head())
```

	Year	Make	Model	Electric Range	Longitude	Lattitute \
0	2020	HYUNDAI	KONA	258	-122.343010	47.659185
1	2022	JEEP	GRAND CHEROKEE	25	-122.205780	47.762405
2	2023	JEEP	GRAND CHEROKEE	25	-120.602720	46.596562
3	2018	TESLA	MODEL 3	215	-122.209285	47.711240
4	2018	BMW	I3	97	-122.896920	47.043535

  

	Price
0	22000.0
1	NaN
2	NaN
3	44000.0
4	NaN

```
In [27]: print(len(df_merged[df_merged["Price"].notna()]))

19745
```

## Problem 4 (10 pts)

- Using the DataFrame from the previous problem, remove records where the **Year** column is for 2015 or earlier. Apply the format `dfmerged = dfmerged[condition]`, choosing the appropriate condition.
- Generate the table, a screenshot of which is provided below, using the `pivot_table()` method and the aggregation function `size`. The entries in the table will indicate the number of cars with the specified make, model, and year in the dataset.

		Year	2016	2017	2018	2019	2020	2021	2022	2023	2024
Make	Model										
ALFA ROMEO	TONALE		0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	12.0
	AUDI	A3	212.0	189.0	173.0	0.0	0.0	0.0	0.0	0.0	0.0
		A7	0.0	0.0	0.0	0.0	0.0	12.0	0.0	0.0	0.0
		A8 E	0.0	0.0	0.0	0.0	3.0	0.0	0.0	0.0	0.0
		E-TRON	0.0	0.0	0.0	443.0	0.0	183.0	228.0	125.0	0.0
		E-TRON GT	0.0	0.0	0.0	0.0	0.0	80.0	41.0	0.0	0.0
		E-TRON SPORTBACK	0.0	0.0	0.0	0.0	26.0	73.0	66.0	0.0	0.0
		Q4	0.0	0.0	0.0	0.0	0.0	82.0	208.0	0.0	0.0
		Q5	0.0	0.0	0.0	0.0	0.0	140.0	0.0	0.0	0.0
		Q5 E	0.0	0.0	0.0	0.0	196.0	283.0	0.0	136.0	0.0
		Q8	0.0	0.0	0.0	0.0	0.0	0.0	0.0	80.0	0.0
		RS E-TRON GT	0.0	0.0	0.0	0.0	0.0	18.0	9.0	0.0	0.0
BENTLEY	BENTAYGA		0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
	FLYING SPUR		0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
BMW	330E		10.0	97.0	84.0	0.0	0.0	64.0	61.0	111.0	0.0
	530E		0.0	0.0	201.0	84.0	16.0	21.0	27.0	59.0	0.0
	740E		0.0	6.0	19.0	5.0	0.0	0.0	0.0	0.0	0.0
	745E		0.0	0.0	0.0	0.0	7.0	0.0	0.0	0.0	0.0
	745LE		0.0	0.0	0.0	0.0	0.0	0.0	2.0	0.0	0.0
	I3		184.0	392.0	239.0	201.0	55.0	44.0	0.0	0.0	0.0
	I4		0.0	0.0	0.0	0.0	0.0	108.0	605.0	0.0	0.0

- Use the `groupby()` method to create a table similar to the one above but this time entries will indicate the average latitude of the car with the specified make, model and year.

```
In [29]: df_merged = df_merged[ df_merged.Year > 2015 ]
```

```
In [31]: df_pivot = df_merged.pivot_table(index=["Make", "Model"], columns="
Year", aggfunc="size", fill_value=0.0)

display(df_pivot.head(20))
```

		Year	2016	2017	2018	2019	2020	2021	2022	2023	2024
Make	Model										
ALFA ROMEO	TONALE	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	12.0
AUDI	A3	212.0	189.0	173.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
	A7	0.0	0.0	0.0	0.0	0.0	0.0	12.0	0.0	0.0	0.0
	A8 E	0.0	0.0	0.0	0.0	3.0	0.0	0.0	0.0	0.0	0.0
	E-TRON	0.0	0.0	0.0	443.0	0.0	183.0	228.0	125.0	0.0	0.0
	E-TRON GT	0.0	0.0	0.0	0.0	0.0	0.0	80.0	41.0	0.0	0.0
	E-TRON SPORTBACK	0.0	0.0	0.0	0.0	26.0	73.0	66.0	0.0	0.0	0.0
	Q4	0.0	0.0	0.0	0.0	0.0	0.0	82.0	208.0	0.0	0.0
	Q5	0.0	0.0	0.0	0.0	0.0	0.0	140.0	0.0	0.0	0.0
	Q5 E	0.0	0.0	0.0	0.0	196.0	283.0	0.0	136.0	0.0	0.0
	Q8	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	80.0
	RS E-TRON GT	0.0	0.0	0.0	0.0	0.0	0.0	18.0	9.0	0.0	0.0
BENTLEY	BENTAYGA	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
	FLYING SPUR	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
BMW	330E	10.0	97.0	84.0	0.0	0.0	64.0	61.0	111.0	0.0	0.0
	530E	0.0	0.0	201.0	84.0	16.0	21.0	27.0	59.0	0.0	0.0
	740E	0.0	6.0	19.0	5.0	0.0	0.0	0.0	0.0	0.0	0.0
	745E	0.0	0.0	0.0	0.0	7.0	0.0	0.0	0.0	0.0	0.0
	745LE	0.0	0.0	0.0	0.0	0.0	0.0	2.0	0.0	0.0	0.0
	I3	184.0	392.0	239.0	201.0	55.0	44.0	0.0	0.0	0.0	0.0

```
In [33]: df_groupby = df_merged[["Make", "Model", "Year", "Latitude"]].groupby(
["Make", "Model", "Year"]).mean("Latitude")

display(df_groupby.head(20))
```

			Latitude
Make	Model	Year	
ALFA ROMEO	TONALE	2024	47.607409
AUDI	A3	2016	47.619987
		2017	47.481332
		2018	47.496019
	A7	2021	47.491215
	A8 E	2020	47.742747
	E-TRON	2019	47.481997
		2021	47.510774
		2022	47.525202
		2023	47.166119
	E-TRON GT	2022	47.586789
		2023	47.359299
E-TRON SPORTBACK		2020	47.543414
		2021	47.440235
		2022	47.463412
	Q4	2022	47.535088
		2023	47.483677
	Q5	2022	47.497235
	Q5 E	2020	47.457532
		2021	47.492323

## Problem 5 (10 pts)

- There is a 3-row, 7-columns table whose code is given below. Use `melt()` to convert that table into this form:

	student_id	Subject_Quarter	Score
0	1	Math_Q1	85
1	2	Math_Q1	90
2	3	Math_Q1	82
3	1	Math_Q2	88
4	2	Math_Q2	85
5	3	Math_Q2	80
6	1	Math_Q3	87
7	2	Math_Q3	83
8	3	Math_Q3	84
9	1	Science_Q1	78
10	2	Science_Q1	88
11	3	Science_Q1	80

```
In [35]: # Sample dataset
data = {
    'student_id': [1, 2, 3],
    'Math_Q1': [85, 90, 82],
    'Math_Q2': [88, 85, 80],
    'Math_Q3': [87, 83, 84],
    'Science_Q1': [78, 88, 80],
    'Science_Q2': [82, 85, 78],
    'Science_Q3': [84, 87, 83]
}
scores_df = pd.DataFrame(data)
scores_df
```

Out[35]:

	student_id	Math_Q1	Math_Q2	Math_Q3	Science_Q1	Science_Q2	Science_Q3
0	1	85	88	87	78	82	84
1	2	90	85	83	88	85	87
2	3	82	80	84	80	78	83

```
In [37]: # your solution goes here
melted_df = pd.melt(scores_df, id_vars=['student_id'], var_name='Subject_Quarter', value_name='Score')

print("\nMelted Data:")
display(melted_df)
```

Melted Data:



	student_id	Subject_Quarter	Score
0	1	Math_Q1	85
1	2	Math_Q1	90
2	3	Math_Q1	82
3	1	Math_Q2	88
4	2	Math_Q2	85
5	3	Math_Q2	80
6	1	Math_Q3	87
7	2	Math_Q3	83
8	3	Math_Q3	84
9	1	Science_Q1	78
10	2	Science_Q1	88
11	3	Science_Q1	80
12	1	Science_Q2	82
13	2	Science_Q2	85
14	3	Science_Q2	78
15	1	Science_Q3	84
16	2	Science_Q3	87
17	3	Science_Q3	83

## Problem 5 - Quality Control in a Manufacturing Plant (10 pts)

Imagine you work as a quality control analyst in a manufacturing plant that produces ball bearings. Each day, multiple batches of ball bearings are produced. To ensure the consistency and quality of the ball bearings, samples from each batch are measured to determine their diameters.

Over the course of a month, you've collected diameter data for these samples from various batches. The objective is to determine the batch consistency by measuring the standard deviation of the diameters. A lower standard deviation would indicate that the ball bearings in a batch are more consistent in size.

- Load the **ball\_bearings.csv** file into a dataframe.
- Use `groupby()` to calculate the standard deviation for each batch.
- Sort the results in descending order wrt standard deviation, showing the batch with highest standard deviation at the top.

```
In [39]: df_ball_bearings = pd.read_csv("ball_bearings.csv")

display(df_ball_bearings.head())
```

	batch_id	diameter
0	1	50.248357
1	1	49.930868
2	1	50.323844
3	1	50.761515
4	1	49.882923

```
In [41]: df_grouped = df_ball_bearings.groupby(["batch_id"])["diameter"].std()

display(df_grouped)
```

batch_id	
1	0.480014
2	0.484019
3	0.410424
4	0.556044
5	0.345405
6	0.511339
7	0.534851
8	0.451574
9	0.502297
10	0.367445
11	0.538856
12	0.593042
13	0.516218
14	0.559482
15	0.474119
16	0.319080
17	0.431002
18	0.362575
19	0.441640
20	0.577886
21	0.446483
22	0.645821
23	0.386267
24	0.608465
25	0.508577
26	0.409027
27	0.425961
28	0.435184
29	0.610599
30	0.388254

Name: diameter, dtype: float64

```
In [43]: display(df_grouped.sort_values(ascending=False))
```

```
batch_id
22      0.645821
29      0.610599
24      0.608465
12      0.593042
20      0.577886
14      0.559482
4       0.556044
11      0.538856
7       0.534851
13      0.516218
6       0.511339
25      0.508577
9       0.502297
2       0.484019
1       0.480014
15      0.474119
8       0.451574
21      0.446483
19      0.441640
28      0.435184
17      0.431002
27      0.425961
3       0.410424
26      0.409027
30      0.388254
23      0.386267
10      0.367445
18      0.362575
5       0.345405
16      0.319080
Name: diameter, dtype: float64
```