

Assignment 3 - Deadline: Oct 16, 2024, Wed 11pm

DSAI 510 Fall 2024

Complete the assignment below and upload both the .ipynb file and its pdf to <https://moodle.bogazici.edu.tr> by the deadline given above. The submission page on Moodle will close automatically after this date and time.

To make a pdf, this may work: Hit CMD+P or CTRL+P, and save it as PDF. You may also use other options from the File menu.

```
# Run this cell first

import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

# Set the display option to show all rows scrolling with a slider
pd.set_option('display.max_rows', None)
# To disable this, run the line below:
# pd.reset_option('display.max_rows')
```

Note:

In the problems below, if they ask "show the number of records that are nonzero", the answer is a number; so you don't need to show the records themselves. But if it asks, "show the records with NaN", it wants you to print those records (rows) containing NaN and other entries, not asking how many such records there are. So be careful about what you're asked.

Problem 1 (10 pts)

- (a) Load the Ames house dataset from the file **train.csv**.
- (b) Display the records with sale price greater than 500000 USD and LotFrontage less than 150 feet. Show only these columns: **Id**, **LotFrontage** and **SalePrice**.
- (c) Print the list of all possible distinct values for the column **SaleCondition** for the records where sale price is greater than 500000 USD and LotFrontage is less than 150 feet.
- (d) Create an interactive scatter plot of LotFrontage versus SalePrice, displaying only the records identified in the previous step. When hovering over the dots, the plot should display the **SaleCondition** in addition to **LotFrontage** and **SalePrice** information.

```
# Break your computations into multiple cells

# a
df = pd.read_csv("train.csv")
```

```
# b
```

```
cols = ["Id", "LotFrontage", "SalePrice"]
```

```
condition1 = df["SalePrice"] > 500000
```

```
condition2 = df["LotFrontage"] < 150
```

```
result_df = df[cols][condition2 & condition1]  
display( result_df )
```

	Id	LotFrontage	SalePrice
178	179	63.0	501837
440	441	105.0	555000
691	692	104.0	755000
769	770	47.0	538000
803	804	107.0	582933
898	899	100.0	611657
1046	1047	85.0	556581
1169	1170	118.0	625000

```
# c
```

```
print(df[condition1 & condition2]["SaleCondition"].unique())
```

```
['Partial' 'Normal']
```


```
# d
```

```
import plotly.express as px
```

```
fig = px.scatter(  
    df[condition1 & condition2],  
    x="LotFrontage",  
    y="SalePrice",  
    hover_name="Id", # Display Id on hover  
    hover_data={  
        "SaleCondition": True # Include SaleCondition in hover data  
    },  
    title="Lot Frontage vs Sale Price",  
)
```

```
# Show the plot
```

```
fig.show()
```



Problem 2 (10 pts)

- Display the list of neighborhood names and 'mean sale price' for those neighborhoods for the records whose SaleCondition is 'Normal'.
- Display the list of neighborhood names and the difference "max sale price - mean sale price" for each neighborhood. (Here "-" is for subtraction.)
- Recreate the boxplot comparing Neighborhood to SalePrice that we made in class. This time, order the neighborhoods based on their medians in ascending order, from left to right. In other words, the neighborhood with the lowest SalePrice median should be on the far left.

```
# a
df_mean_neigh = df[["Neighborhood", "SalePrice"]][df.SaleCondition ==
"Normal"].groupby("Neighborhood").mean("SalePrice")
display(df_mean_neigh)
```

Neighborhood	SalePrice
Blmngtn	188977.083333
Blueste	137500.000000
BrDale	107916.666667
BrkSide	125588.425926
ClearCr	220993.000000
CollgCr	193877.224806
Crawfor	204863.651163
Edwards	127803.048780
Gilbert	189392.812500
IDOTRR	108575.862069
MeadowV	98987.500000
Mitchel	155410.714286
NAmes	147533.550505
NPkVill	143031.250000
NWAmes	193799.296875
NoRidge	328219.135135

NridgHt	285046.666667
OldTown	133173.489362
SWISU	139788.636364
Sawyer	136976.611940
SawyerW	191505.600000
Somerst	217760.653061
StoneBr	264870.375000
Timber	238484.392857
Veenker	238772.727273

b

```
df_max_neigh = df[["Neighborhood", "SalePrice"]][df.SaleCondition ==
"Normal"].groupby("Neighborhood").max("SalePrice")
```

```
display(df_max_neigh - df_mean_neigh)
```

	SalePrice
Neighborhood	
Blmngtn	45022.916667
Blueste	13500.000000
BrDale	17083.333333
BrkSide	97911.574074
ClearCr	107007.000000
CollgCr	119122.775194
Crawfor	176136.348837
Edwards	192196.951220
Gilbert	188107.187500
IDOTRR	60924.137931
MeadowV	52412.500000
Mitchel	115589.285714
NAmes	197466.449495
NPkVill	11968.750000
NWAmes	106000.703125
NoRidge	426780.864865
NridgHt	269953.333333
OldTown	341826.510638
SWISU	57211.363636
Sawyer	53023.388060
SawyerW	128494.400000
Somerst	122239.346939
StoneBr	273129.625000
Timber	140015.607143
Veenker	146227.272727

create the index based on the description of question

```
df_index=pd.DataFrame(
    df.groupby(['Neighborhood'])['SalePrice']
    .median())
```

```


.sort_values(ascending=True)).index

# Create a boxplot ordered by neighborhood median sale prices
plt.figure(figsize=(12, 8))
sns.boxplot(x='Neighborhood', y='SalePrice', data=df, order=df_index,
            showmeans=True,
            meanprops={"marker": "o", "markerfacecolor": "white",
                       "markeredgecolor": "black"}, whis=1.5)

# Rotate on x-axis
plt.xticks(rotation=90)

# Show the plot
plt.tight_layout()
plt.show()

```



Problem 3 (10 pts)


Here we'll show some of the houses on the map.

(a) Install the folium package with `!pip install folium`.

(b) Suppose your manager at the yellow website fromtheowner.com wants you to make a webpage showing houses on sale. Create the interactive map by using `folium` package to produce the map shown below for the 30 houses whose Id's and coordinates are given in

locations.csv. When you click on any pin on the map, the box should show the Id and SalePrice of that house as shown in the map below. You can find the SalePrice information in **train.csv**, and it's connected to **locations.csv** by the common column **Id**.

(Hint for folium usage: ChatGPT, Google, folium documentary...)



```
# !pip install folium

df_location = pd.read_csv("locations.csv")
df_train = pd.read_csv("train.csv")

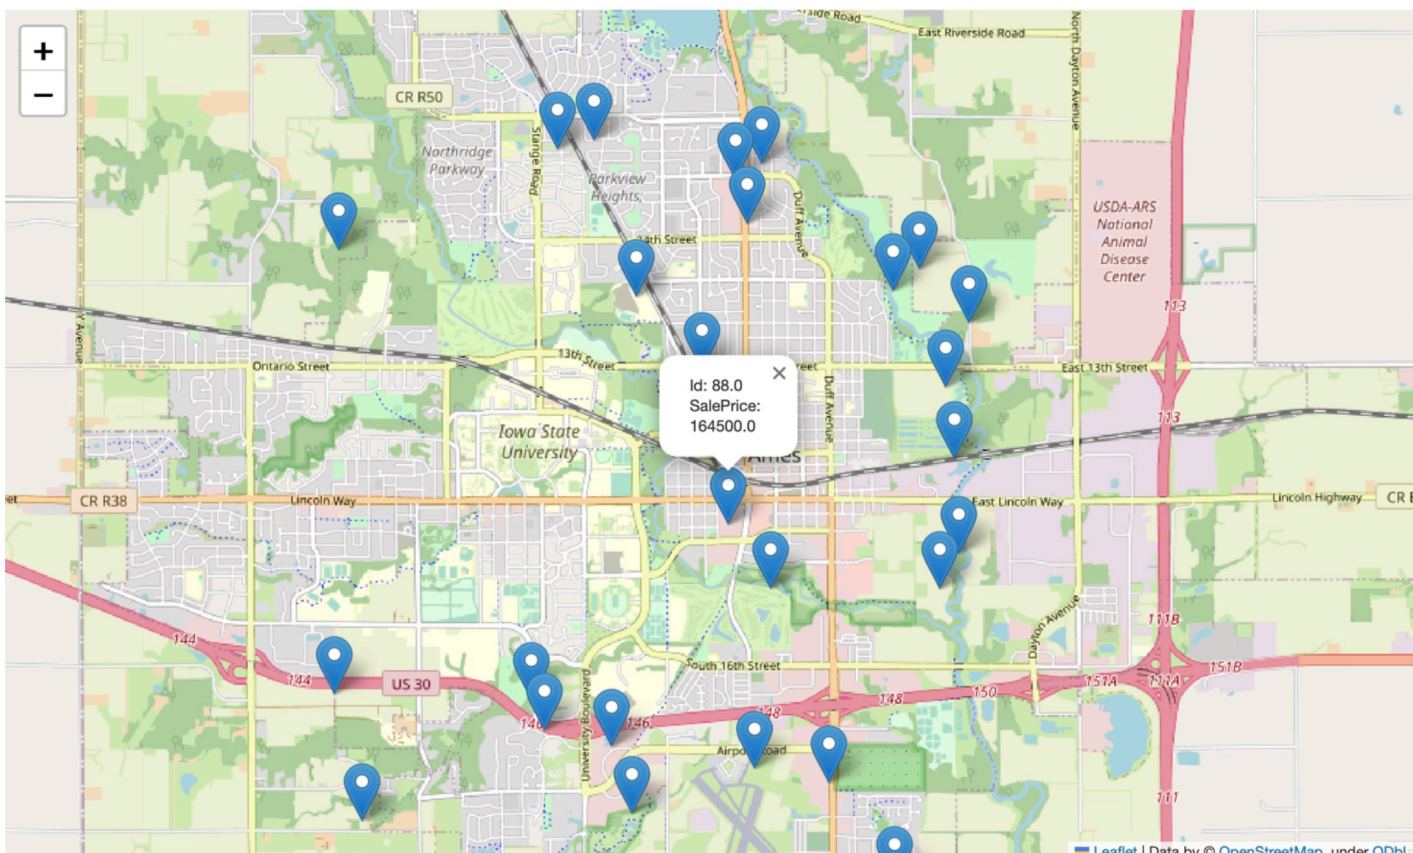
# Merge two data on "Id" using left join
df_merged = df_location.merge(df_train[["Id", "SalePrice"]], on="Id",
                               how="left", suffixes=('_location', '_train'))

import folium

# Step 1: Create a map centered at a specific latitude and longitude
coordinates = list(df_merged.mean())
map_center = coordinates[1:3] # Average of the latitude and longitude
for our dataset
map = folium.Map( location=map_center, zoom_start=13 )

# Step 2: Add markers to the map
def mapper(i):
    point_data = list(i)
    folium.Marker(
        location=[point_data[1], point_data[2]],
        popup=f'Id: {point_data[0]}\nSalePrice: {point_data[3]}',
        icon=folium.Icon(color='blue')
    ).add_to(map)

for i in range(len(df_merged)):
    mapper(df_merged.iloc[i])
```



```
# Step 3: Display the map  
map
```

```
<folium.folium.Map at 0x151f859d0>
```


Problem 4 (10 pts)

(a) Load the data from **Electric_Vehicle_Population_Data.csv** into a dataframe **df**. Show the first five records (don't run **df** to show all records; jupyter notebook crashes as the data has 150482 rows).

(b) Make **df2** where it only includes the records whose **Electric Range** is 'Battery Electric Vehicle (BEV)' and **Electric Range** is greater than zero. (There should be ~47000 records satisfying these conditions; check the length of your final dataframe before proceeding!).

(c) Use **df2** to plot the histogram of the column **Electric Range**

(d) Use **df2** to create three boxplots as we did in the class for Electric Range in the y-axis and 'Model Year', 'Make' and 'Model' categories in the x-axis. Use sns library and set **col_wrap=1**, **sharex=False**, **sharey=False**, **height=10** so that we don't get two or more boxplots side by side. Your plots should look like this (I put them side by side to save space here; yours will be stacked vertically in the Jupyter notebook):



(e) "What story do these boxplots convey?" To answer this question, write at least two observations for each of the three boxplots (in total at least six observations).

(f) Based on the box plots, does any of 'Model Year', 'Make' and 'Model' not determine the **Electric Range**, or do all of these determine it?

```
# a  
  
df = pd.read_csv("Electric_Vehicle_Population_Data.csv")  
  
df.head()
```


	County	Model	Year	Make	Model \
0	King		2020	HYUNDAI	KONA
1	King		2022	JEEP	GRAND CHEROKEE
2	Yakima		2023	JEEP	GRAND CHEROKEE
3	King		2018	TESLA	MODEL 3
4	Thurston		2018	BMW	I3

	Electric Vehicle Type	Electric Range
0	Battery Electric Vehicle (BEV)	258
1	Plug-in Hybrid Electric Vehicle (PHEV)	25
2	Plug-in Hybrid Electric Vehicle (PHEV)	25
3	Battery Electric Vehicle (BEV)	215
4	Plug-in Hybrid Electric Vehicle (PHEV)	97

b

```
df2 = df[(df["Electric Vehicle Type"] == "Battery Electric Vehicle (BEV)") & (df["Electric Range"] > 0)]
```

```
print(len(df2))
```

47109

c

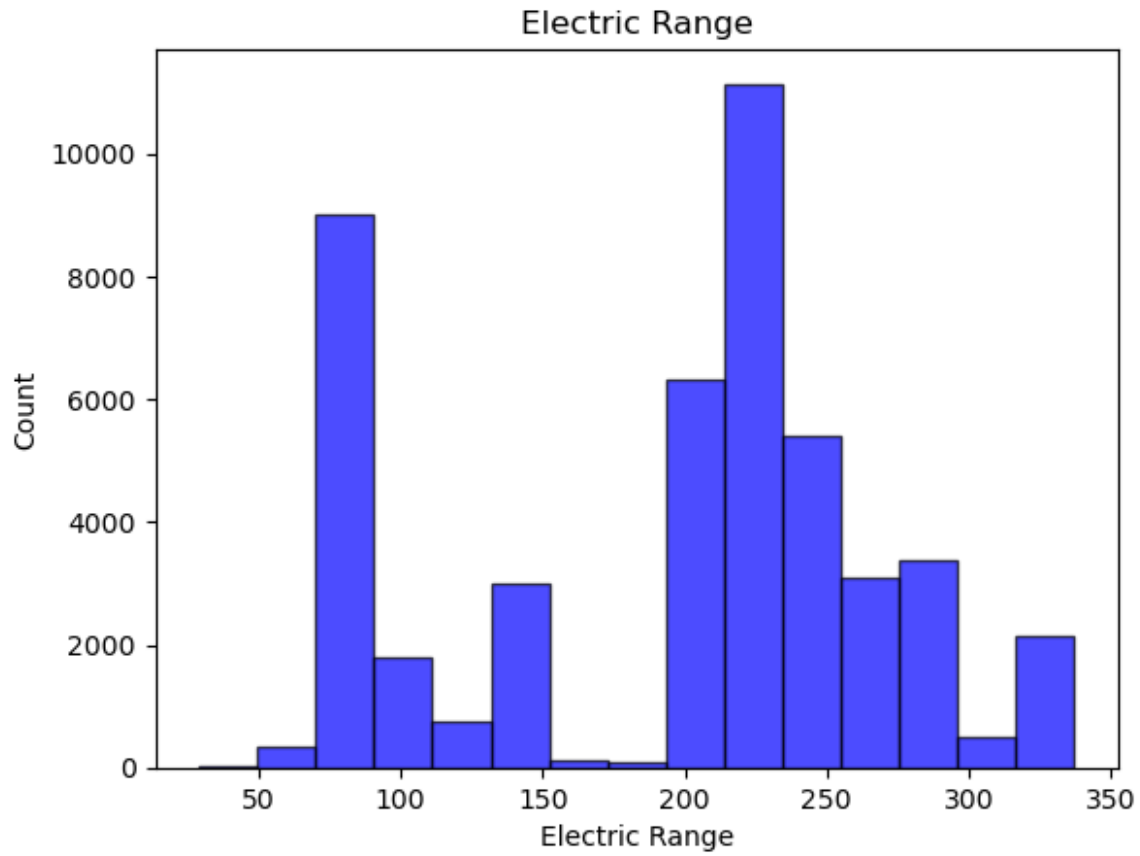
```
import matplotlib.pyplot as plt
```

Histogram for column 'Electric Range'

```
df2["Electric Range"].plot(kind='hist', bins=15, alpha=0.7, color='blue', edgecolor='black')
plt.title('Electric Range')
plt.xlabel('Electric Range')
plt.ylabel('Count')
```

Show the plot

```
plt.show()
```



```
# d



import matplotlib.pyplot as plt
import seaborn as sns

x_ = 'Electric Range'
y_ = ['Model Year', 'Make', 'Model']

# Define boxplot fnc we'll use below
def boxplot(x, y, **kwargs):
    # lines within the boxes will show medians, white dots will show means
    sns.boxplot(x=x, y=y, showmeans=True, meanprops={"marker": "o",
"markerfacecolor": "white", "markeredgecolor": "black"})
    x=plt.xticks(rotation=90)

# Taken from class notes
# It takes original data and melt with x_ and y_ values
melted_categorical = pd.melt(df2, id_vars=x_, value_vars=y_)
g = sns.FacetGrid(melted_categorical, col="variable", col_wrap=1,
```

```
sharex=False, sharey=False, height=10)  
g = g.map(boxplot, "value", x_)
```



```

# e

# Plot 1

# Observation 1
    # Technological advancement is seen clearly, in the first plot, as
    the average range is increased.

# Observation 2
    # There is an instant increase in the difference Q3 - Q1 after
    2013. So, there is an increase on the variety of electric vehicles in
    the market

# Plot 2

# Observation 1
    # There is no corrolation between Make. Some Brands can make
    better cars in terms of range and some do not. Still some are better
    than other like Tesla
# Observation 2
    # The boxplots indicates that the data is left skewed mostly.
    Because brands created middle range vehicles usually.
    # Majority of their work goes to a focused model to increase more
    ranged vehicles.

# Plot 3

# Observation 1
    # If the model is heavy then the range falls like with Ranger or
    Pickup.
# Observation 2
    # TESLA has one range option only in MODEL Y. This is like its
    special car.

# f
    # There is a corrolation between Electric Range with Year and
    Model, yet Make needs more debate. Mean for Make plots does not
    preserve the increase in the Range

```