510 DATA SCIENCE

Lecture 01

Fall 2024

Instructor: Assoc. Prof. Şener Özönder

Email: sener.ozonder@bogazici.edu.tr

Institute for Data Science & Artificial Intelligence

Boğaziçi University

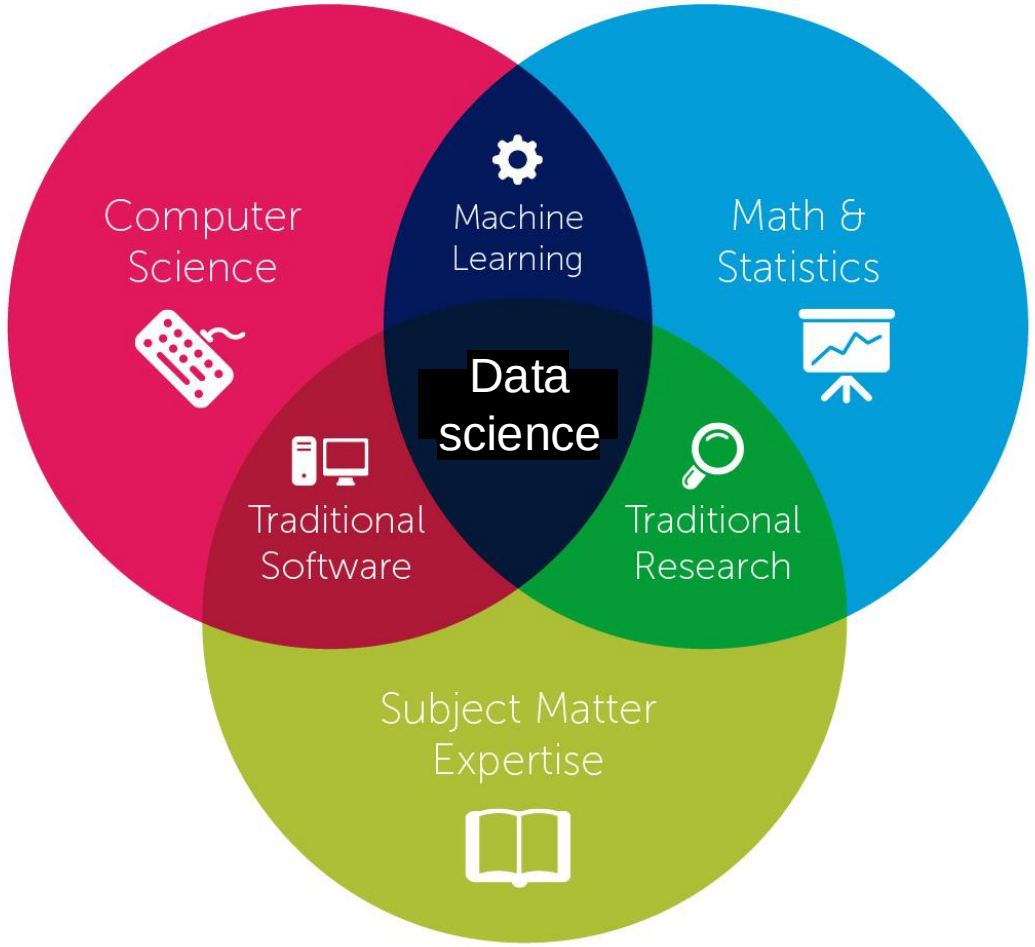

What's in this class?

- In this course, we will focus on core data science concepts such as data cleaning, descriptive statistics, transforming data, modelling (regression, classification, clustering) and prediction/inference. (Check syllabus & lecture slides at http://moodle.bogazici.edu.tr)

Week	Subject
1 (First class: Sept 23)	What is data and data science?
2	Data types, data collection and databases
3	Data wrangling with pandas
4	Exploratory Data Analysis, Feature Engineering and Visualization
5	Probability and Statistics for Data Scientists
6	Modeling: Linear Regression
7	Model evaluation metrics
	Midterm
8	Modeling: Classification
9	Dimensional reduction
10	Modeling: Clustering
11	Modeling: Clustering
12	Time series analysis
13	Processing and modeling text data
14	Best Practices
	Final exam

What's *not* in this class?

- We'll do some of the things below a little bit but these subjects will be treated in depth in their respective courses.

- **No data engineering**

520 Big Data Systems
523 Cloud Computing and Distributed Systems
524 Software Design for Data Science

- **No dashboarding but some viz**

522 Business Intelligence and Analytics
521 Data Visualization for Data Scientists

- **Very little image processing**

543 Image Processing with Machine Learning
544 Computer Vision with Machine Learning

- **A bit text data processing**

545 Natural Language Processing

- **No web scraping**

526 Web Mining
531 Social Media Analytics

- **Some basic machine learning, but not deep learning**

512 Machine Learning
541 Deep Learning (2nd semester)

- **Some statistics, but not in depth.**

514 Statistical Inference

Python *vs* R

- We will use Python because with Python you can not only do data science and machine learning, but also things like web and application development, dashboarding etc.

IDE (Integrated Development Environment) is where you write code. For Python, what's mostly used is Jupyter Notebook (Anaconda). You can also use Google Colab to code in Python even in your cellphone; it's an online IDE and it uses Google's computational resources. Another option is Microsoft's VS Code, where you can code in almost in all languages. Start with Jupyter and Google Colab.

- R is originally developed for statistical analysis. Mostly statisticians use it. Transforming data and statistical tests sometimes easier in R than Python, but it will *not* allow you to build an application to deploy except some web dashboards. Still, if you're collaborating with statisticians or bio-statistics people and if they're already using R, then you can easily learn R code in a week.

To use R, install RStudio IDE: https://posit.co/products/open-source/rstudio/

Python *vs* R

Jupyter Notebook

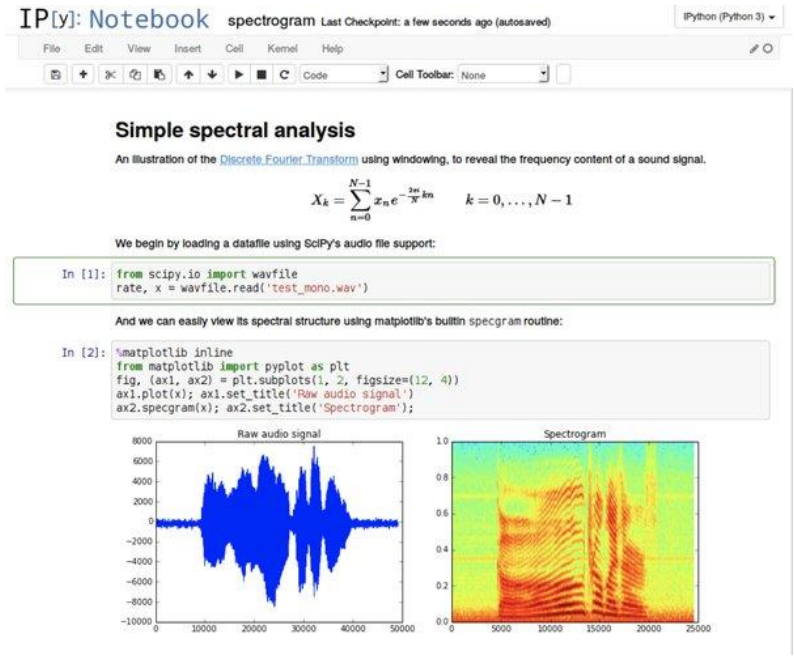

Google Colab

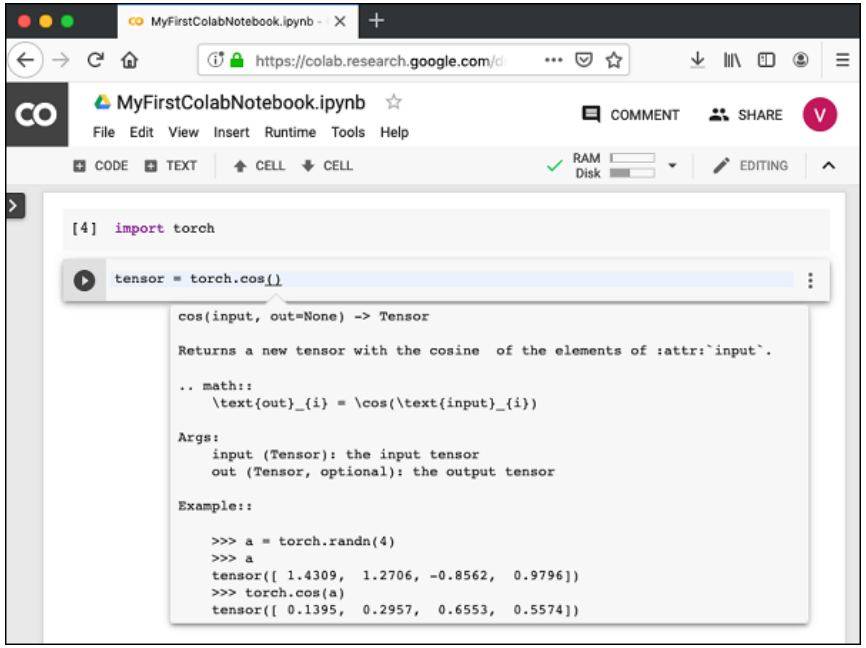

R Studio

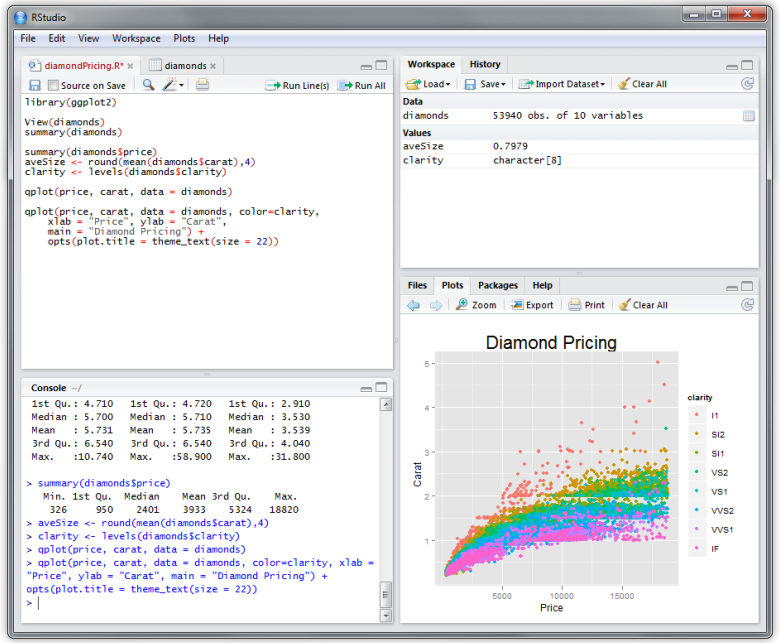

ChatGPT policy

I think it started reasoning in 2024!

- ChatGPT is a 30 year old dillattente or encyclopedist; can't reason much.
- Can you use ChatGPT/Claude/Gemine in this class? Not "you can", "you should." Why look up Google rather than simply asking ChatGPT to find the piece of code you need?
- In the pre-ChatGPT era, people relied on websites like Stack Overflow for codes. Now, ChatGPT can provide the code you need through a simple chat interface.
- But here's the caviat: ChatGPT is great as a conversational assistant or an encyclopedist to lead you to the right direction, but don't rely on it as it may mislead you. So use it with caution!

ChatGPT policy

From the Syllabus:

Collaboration and ChatGPT policy for homeworks: Collaborating with fellow students on assignments is both permitted and encouraged. You're also welcome to use ChatGPT for conceptual clarity and guidance on Python coding. However, all submitted work must fundamentally be your own. Verbatim copying of answers from peers or ChatGPT is prohibited. While you can use code directly from ChatGPT, ensure you understand its structure and logic thoroughly, as the instructor will ask you to explain any code you submit in the midterm and final exams.

What is data science?

- An umbrella term for creating impact for business or research problems by using both statistics/math and computer science together.
- Sometimes "data science" is defined very broadly to include collection, cleaning of data including setting up the hardware and software infrastructure as well as creating dashboards to show the analytics, but these roles have different names.

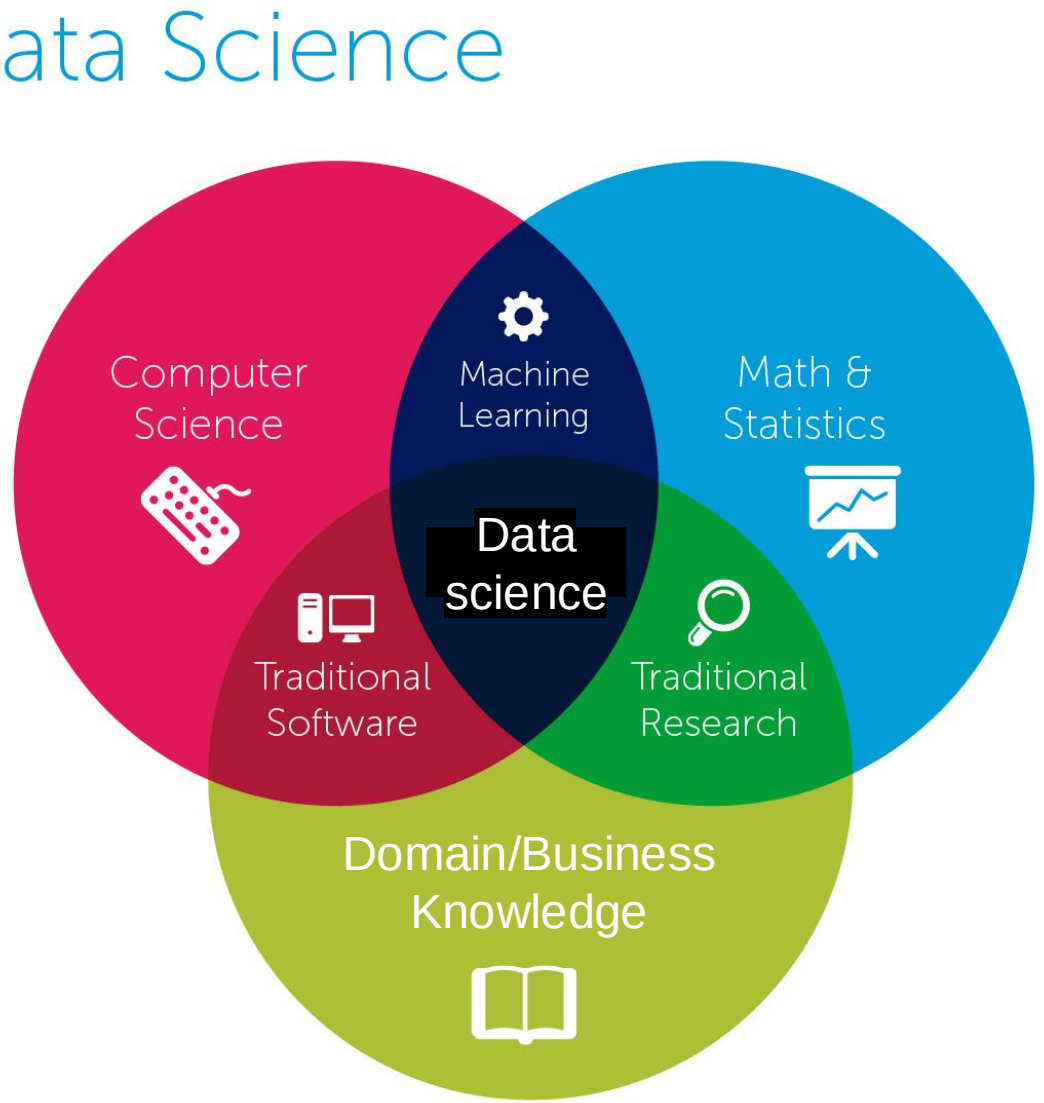

Values created by a data scientist

Some examples:

- **Industry**: Process Optimization, Quality Control, Predictive Maintenance, Energy Efficiency
- **Business**: Customer Insights, Marketing Optimization, Sales Forecasting, Risk Management, Inventory Management, Fraud Detection, Churn Prediction, Pricing Strategy, Operational Efficiency and anything else that will increase revenue.
- **Academia**: Gather insights from data, test hypotheses, discover unknown structures and relations within the data, simulate and model complex systems,

What happened to good old statistics?

- Data science and machine learning is also called “statistical learning”. Data science sometimes has more emphasis on black box modeling such as deep learning of big data for the purpose of prediction/inference where explainability is often sacrificed. Statisticians don’t like this.
- When you need to not just model but understand the data (“data mining”) and extract knowledge and causal relations from data (“knowledge discovery”), good old statistics may become more important than data science.
- Having said that, the intersection of these two approaches grows everyday.

Data Science – Baba Brinkman Music Video

105 B görüntüleme • 3 yıl önce

Baba Brinkman

Rap battle between a **data scientist** and a classical **statistician**, arguing f

0:02 Data Science **Statisticians Data Science Statisticians Data Scienc**

Altyazılar

http://youtu.be/uHGlCi9jOWY

Who is a data scientist?

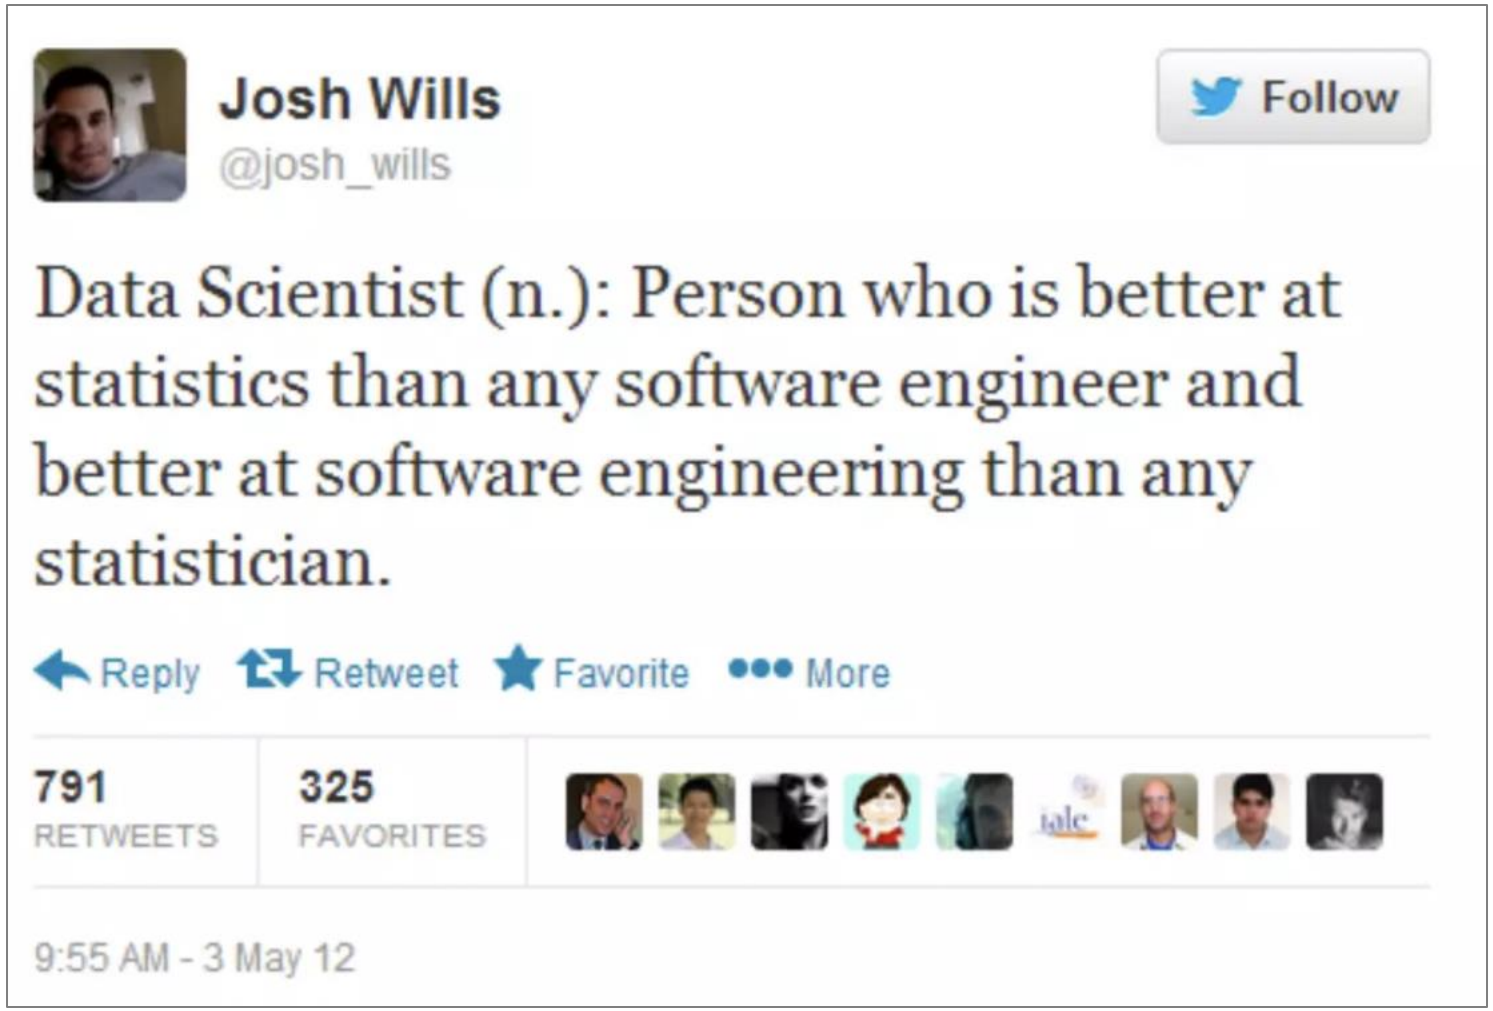

Data science and ML is industry driven!

- Data science and machine learning are industry driven because big data is mostly produced by industry not academia. Also big data is not publicly shared for academicians to play with. (Industry: Hospitals, tech companies, financial sector etc.)
- Machine learning (ML) is not like theoretical physics where you can discover new particles with only math. Most of the time, ML models are designed and optimized specifically for the dataset at hand, so without real data you can't develop good models on fake/synthetic data. *ML models learns from the data!*
- Also developments in computer hardware and software, emergence of cheat IoT sensors, collection of big data by tech companies (Google, Facebook etc.) created the field of data science besides statistics. Data science started to become famous around 2010s.

Data roles

Data Scientist

uses statistics and machine learning to make predictions and answer key business questions

Skills - Math, Programming, Statistics


Tech - SQL, Python, R, Cloud

Data Engineer

build and optimize the systems that allow data scientists and analysts to perform their work

Skills - Programming, BigData & Cloud


Tech - SQL, Python, Cloud, Distributed Computing

Data Analyst

deliver value by taking data, communicating the results to help make business decisions

Skills - Communication, Business Knowledge


Tech - SQL, Excel, Tableau

Data roles

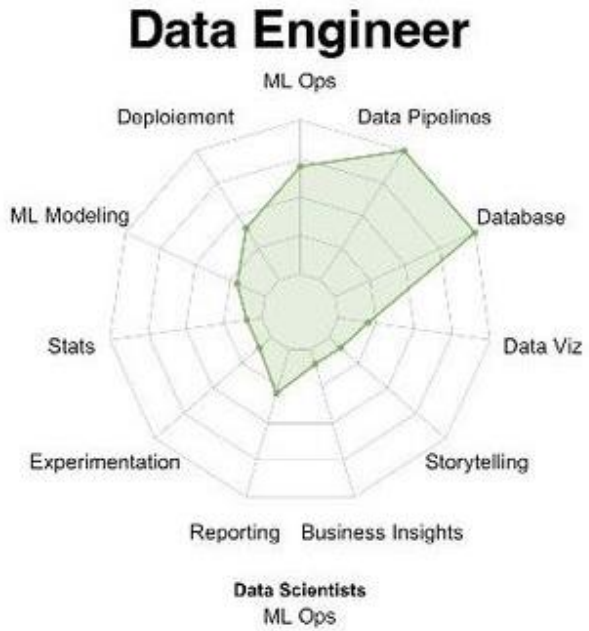
Data Engineer
ML Ops
Data Pipelines
Database
Data Viz
Storytelling
Business Insights
Reporting
Experimentation
Stats
ML Modeling
Deployement
Data Scientists
ML Ops


ML Engineer
ML Ops
Data Pipelines
Database
Data Viz
Storytelling
Business Insights
Reporting
Experimentation
Stats
ML Modeling
Deployement
Data Analysts
ML Ops


Data Scientist
ML Ops
Data Pipelines
Database
Data Viz
Storytelling
Business Insights
Reporting
Experimentation
Stats
ML Modeling
Deployement

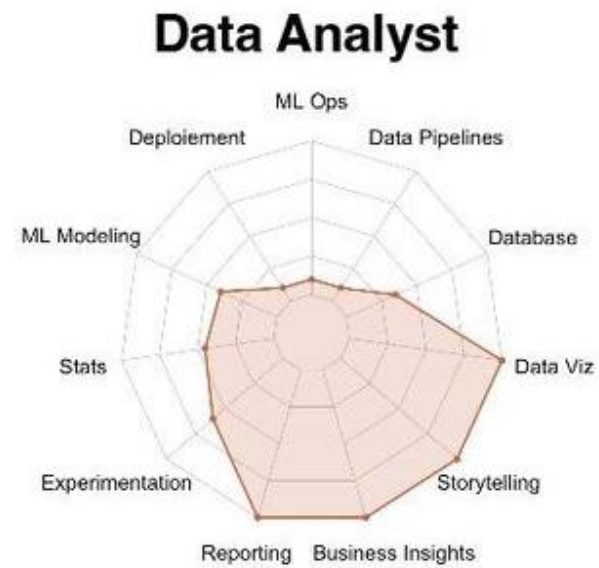
Data Analyst
ML Ops
Data Pipelines
Database
Data Viz
Storytelling
Business Insights
Reporting
Experimentation
Stats
ML Modeling
Deployement

Data roles

- Even though there are overlaps between them, some roles are directly related to some DSAI courses.

(See course descriptions at https://dsai.boun.edu.tr/course-descriptions)

- You can choose to specialize in one role, but it's always good to have basic level skills in all.
- The most important skill is to understand the business/research problem at hand and being able to formulate the solution and determine the necessary tools for it. This is why you're here!

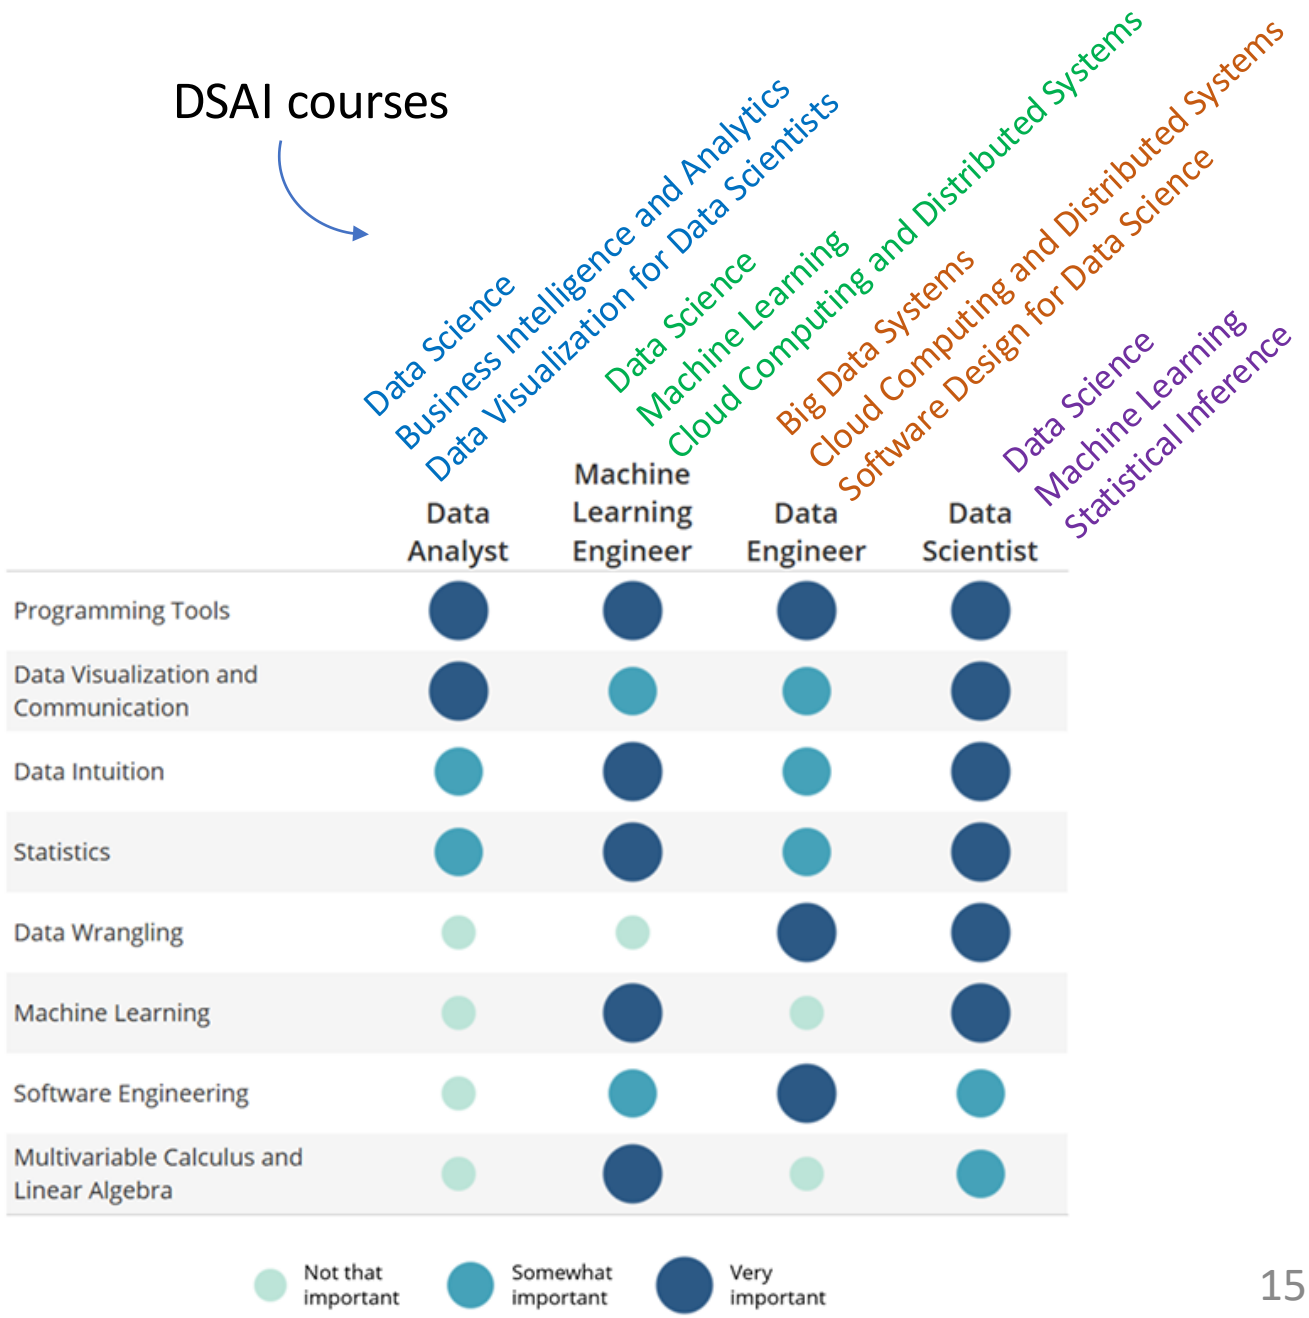

Data roles

- Caution: The definition of these roles are not set in stone yet somewhat company and country dependent and still evolving.
- Always define well what you mean by these words while communicating.

5 Types of Analytics

1. Descriptive: What is happening?

- Correct Data
- Effective Exploratory data analysis

2. Diagnostic: Why is it happening?

- Finding the causes
- Separating all the patterns

3. Predictive: What is likely to happen?

- Choosing the right algorithm
- Bulding the right business strategies

4. Prescriptive: What do I need to do?

- Using the advance analytics
- Recommended actions

5. Cognitive Analytics

- Neurological and Behavioral analysis

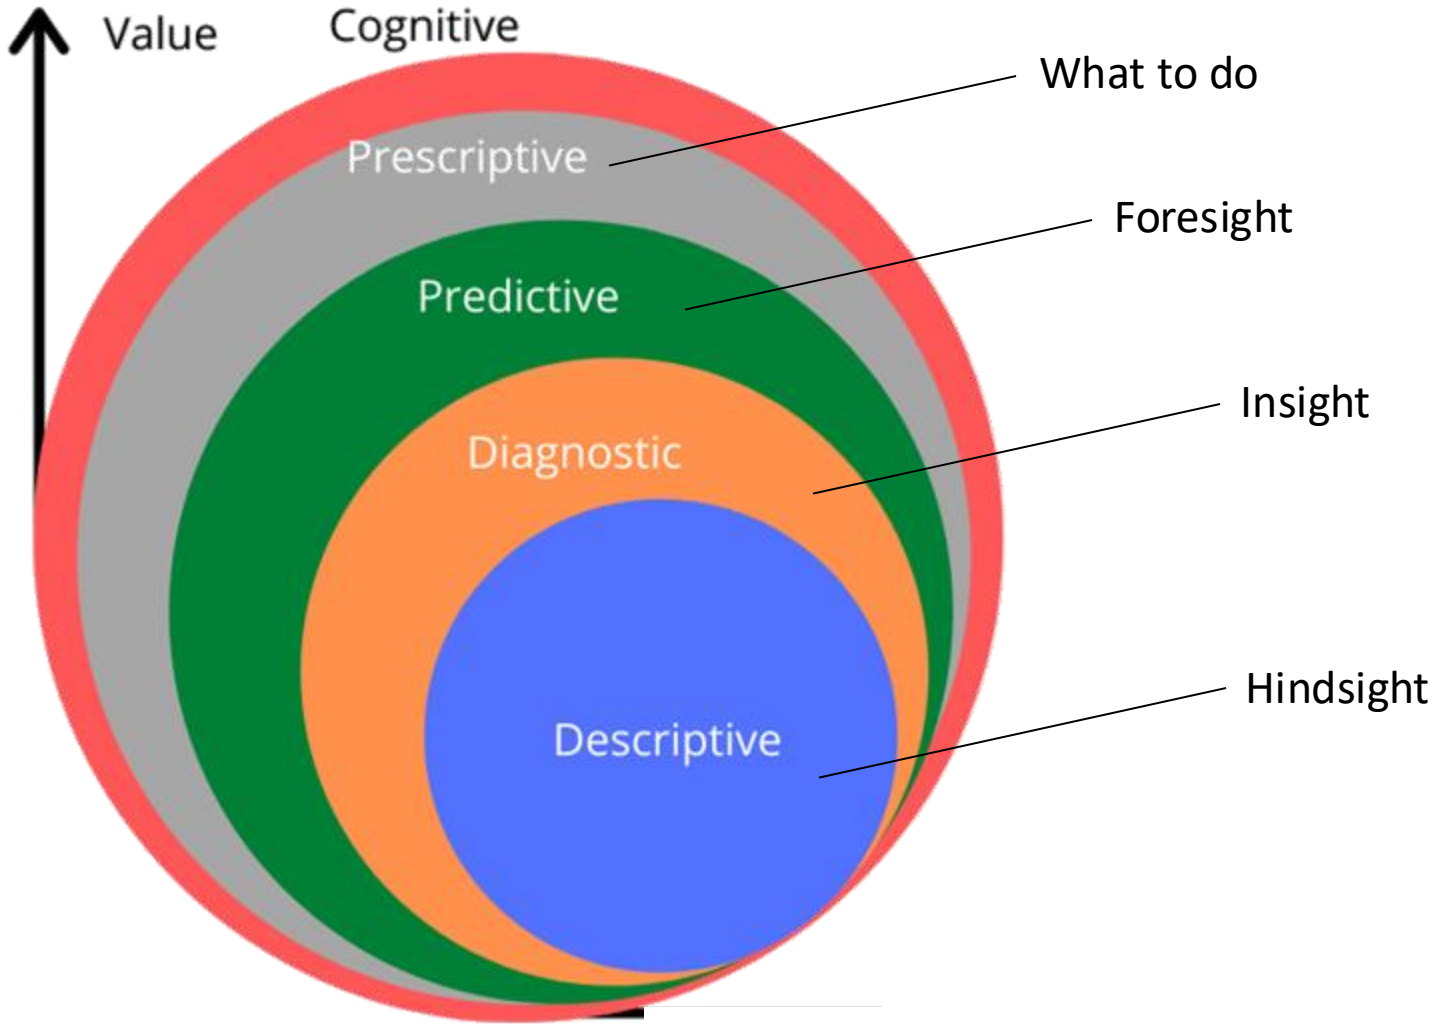

5 V’s of Big data

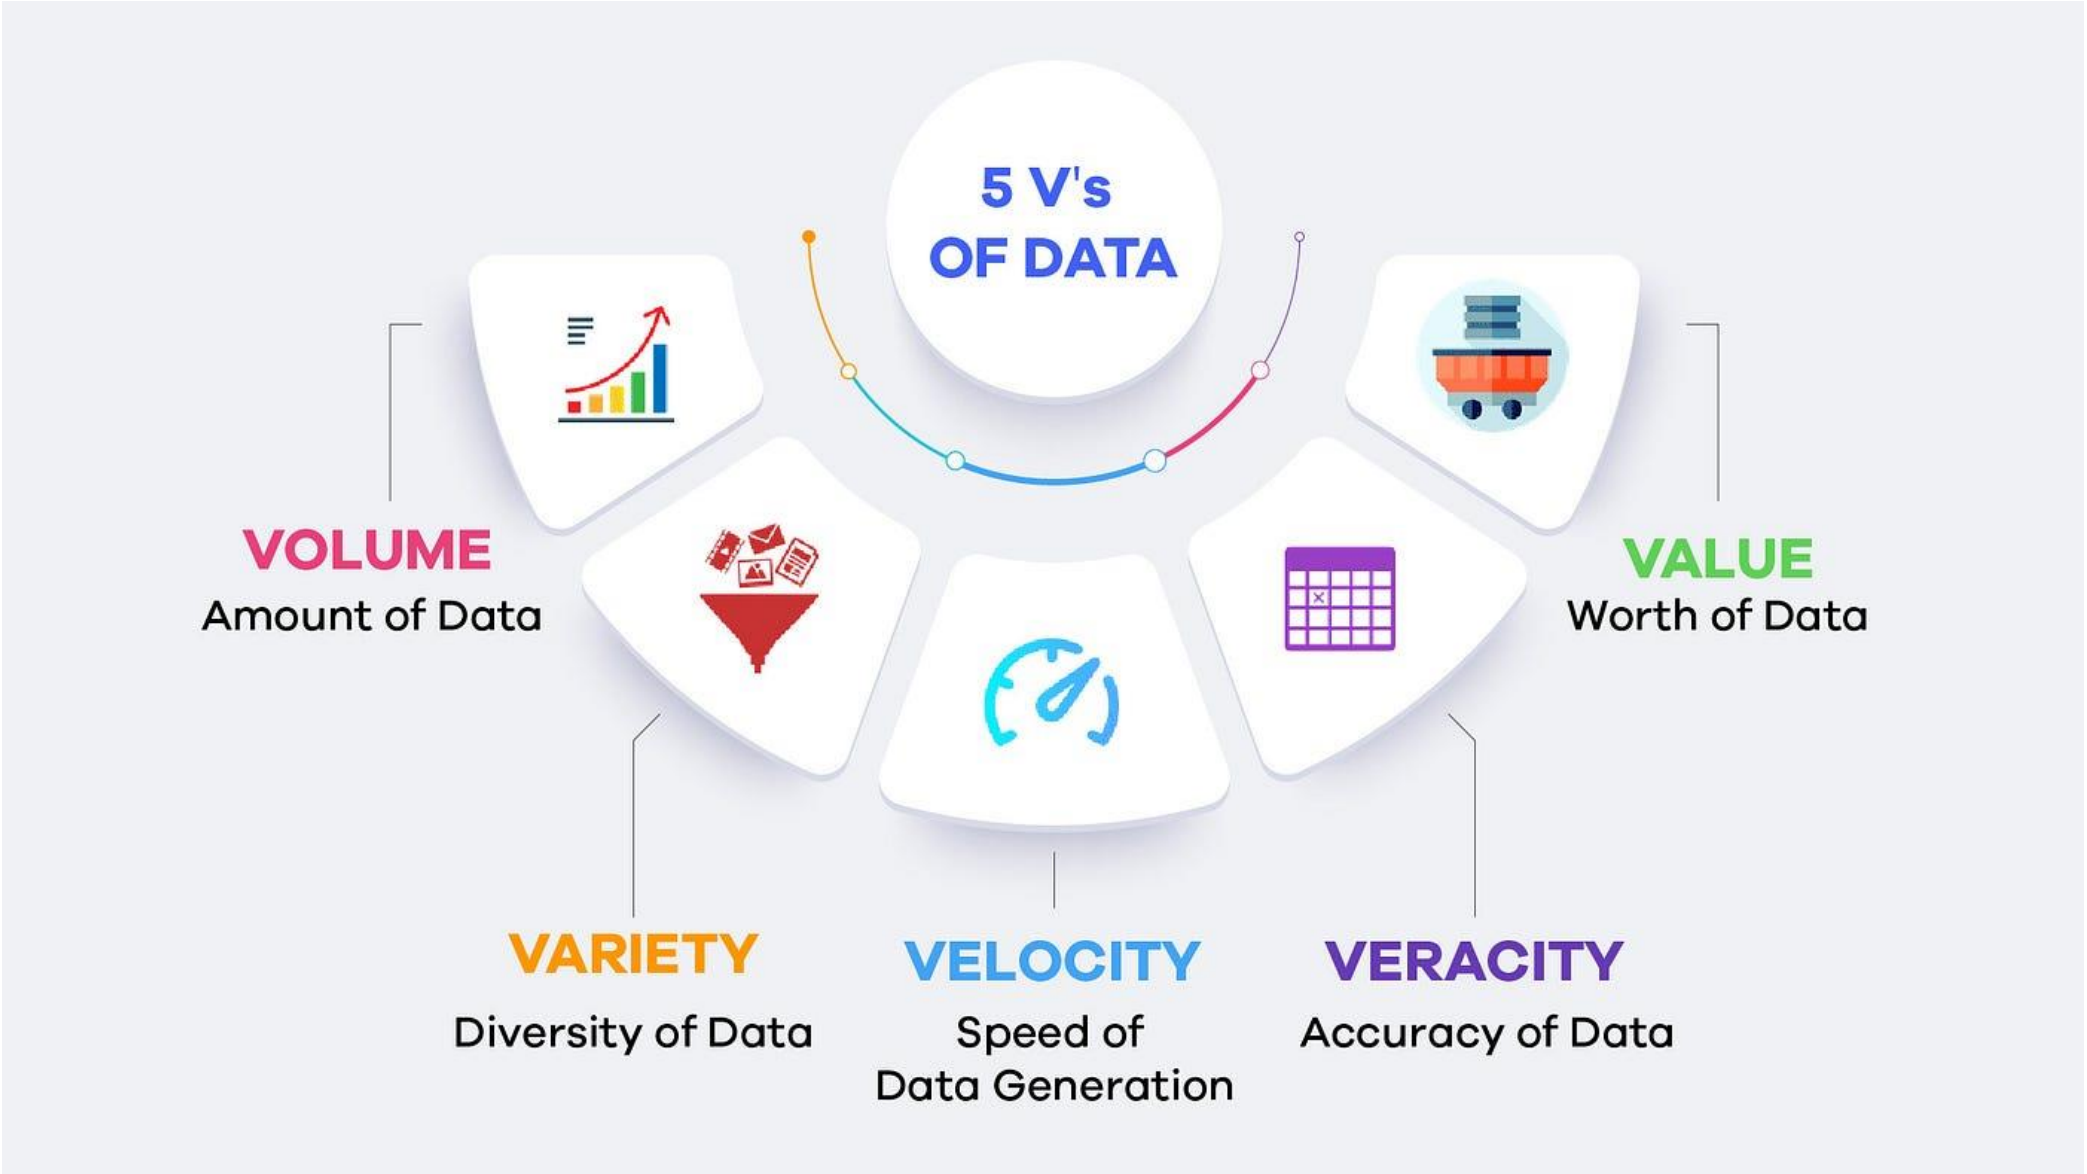

Data Science Lifecycle

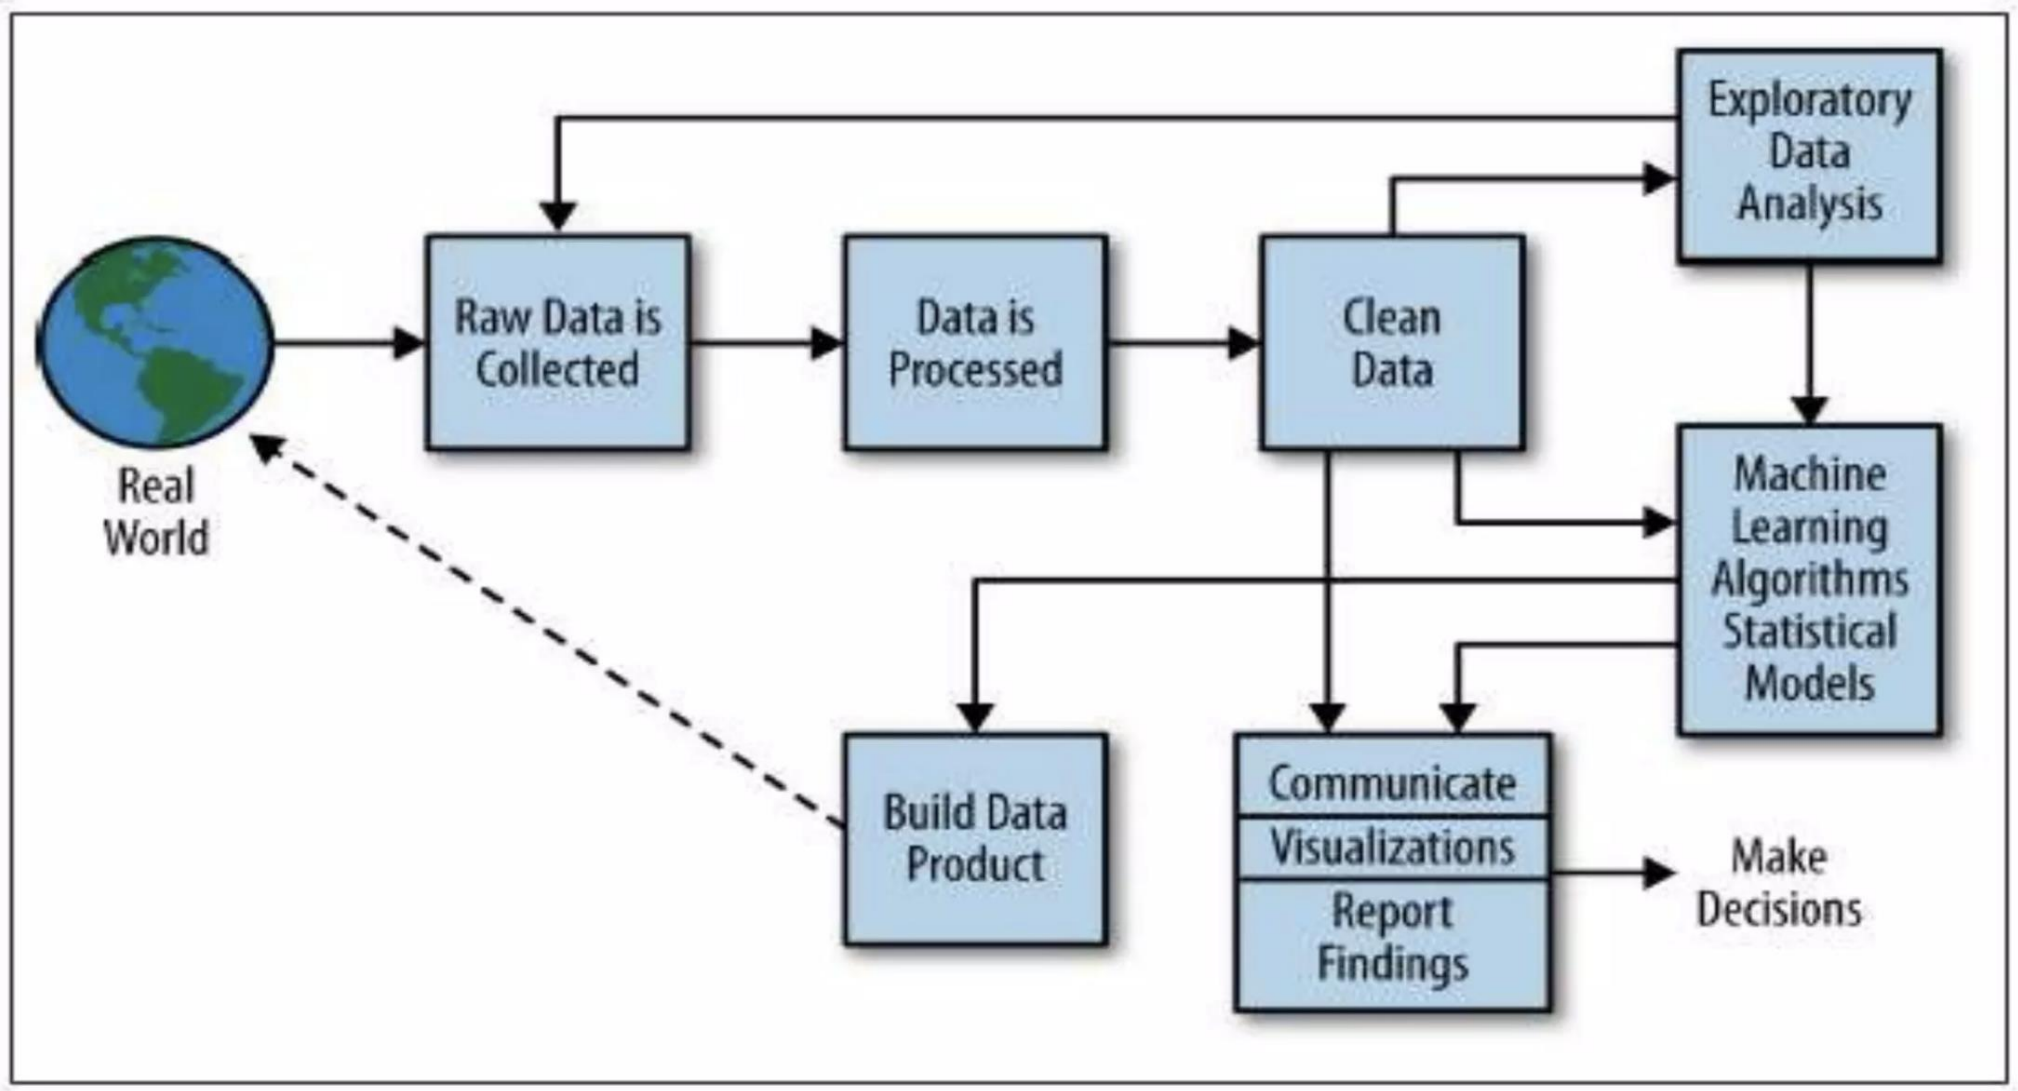

Tabular Data

Column, feature, attribute, variable.

Row, instance, point, vector, record, entry.

Name	Age	IQ	Wealth (k$)
Ismail	30	56	759
Frank	28	117	217
Lucia	99	71	155
Syed	69	104	733

Data lives mathematical spaces

Same box in two different angles:

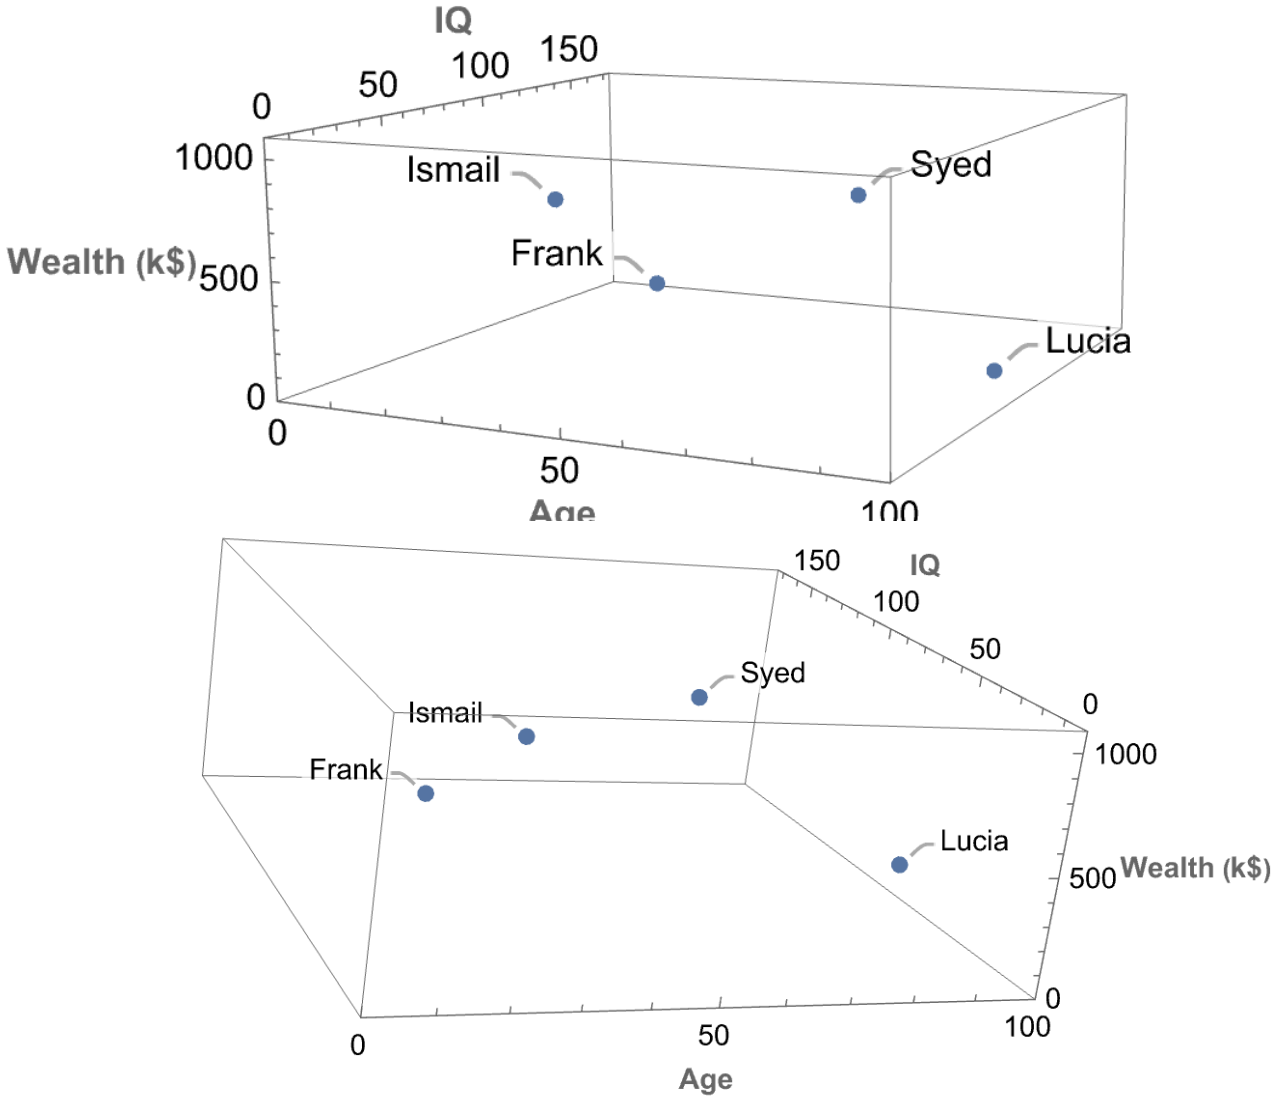

Name	Age	IQ	Wealth (k$)
Ismail	30	56	759
Frank	28	117	217
Lucia	99	71	155
Syed	69	104	733

Tabular Data

Name	Age	IQ	Wealth (k$)
Ismail	30	56	759
Frank	28	117	217
Lucia	99	71	155
Syed	69	104	733

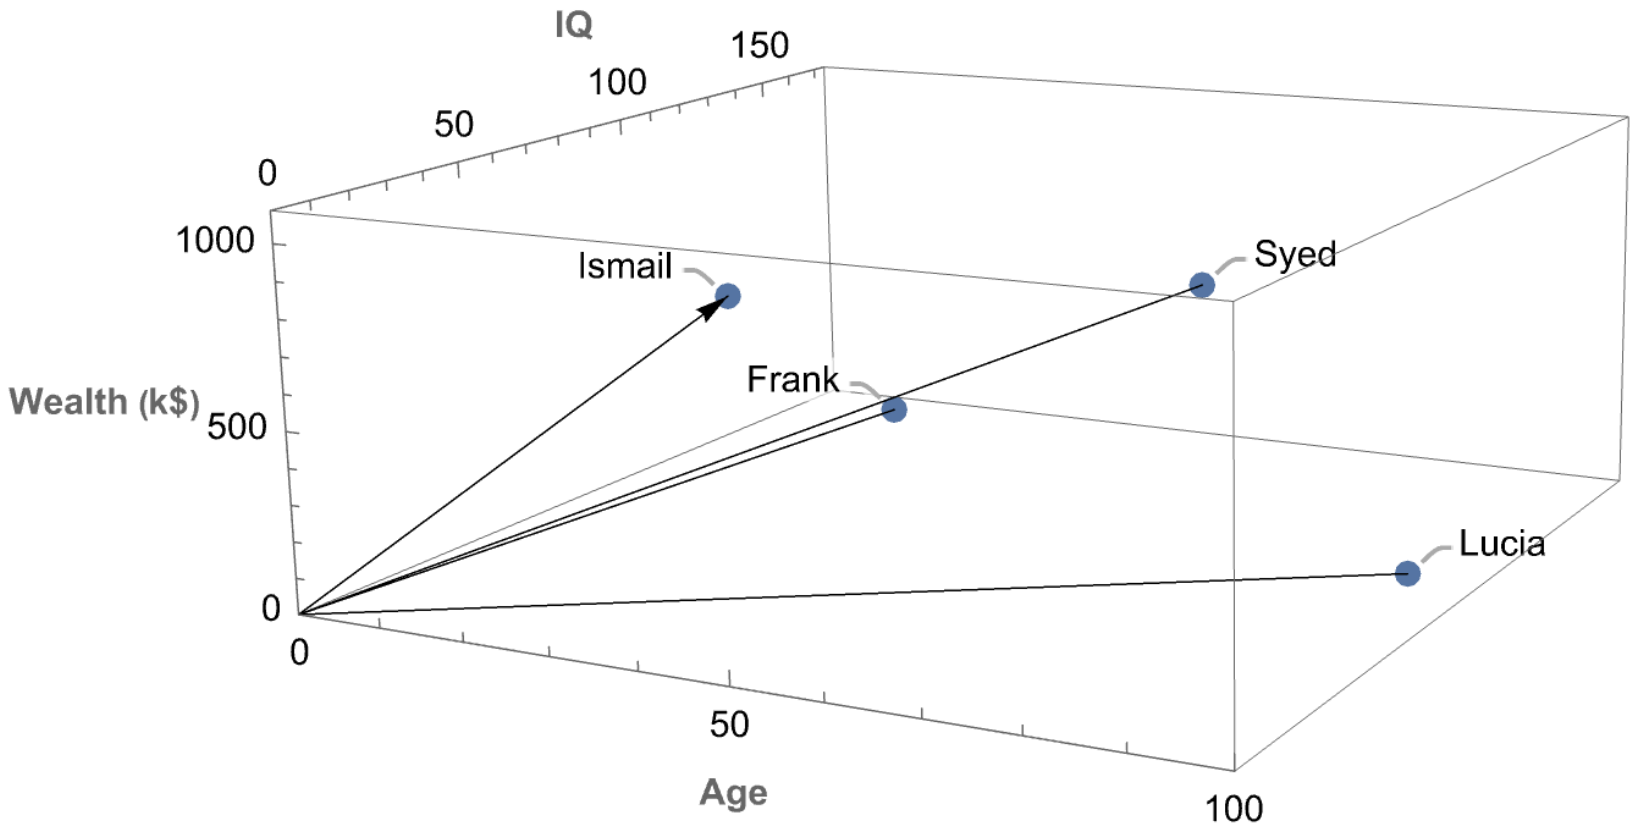

30	56	759

is one data "point" or "vector" in the 3D space where all data points live. In computer, this data point/vector is shown as [30, 56, 759]. You can consider these three numbers as the coordinates in the 3D space of "Age", "IQ" and "Wealth (k$)", or you can show this data as a vector from origin [0, 0, 0] to [30, 56, 759].

Tabular Data

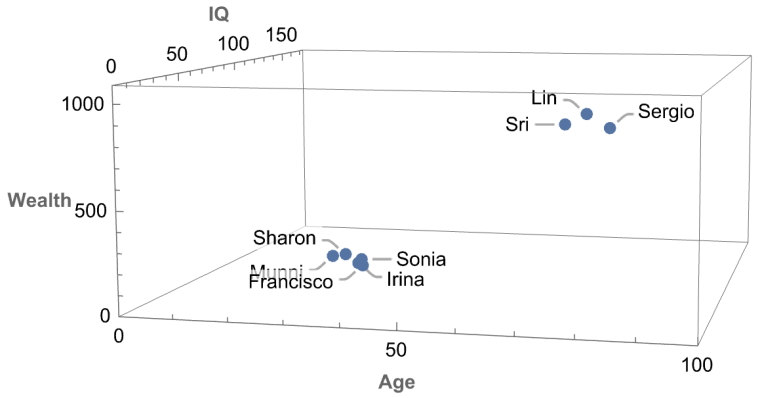

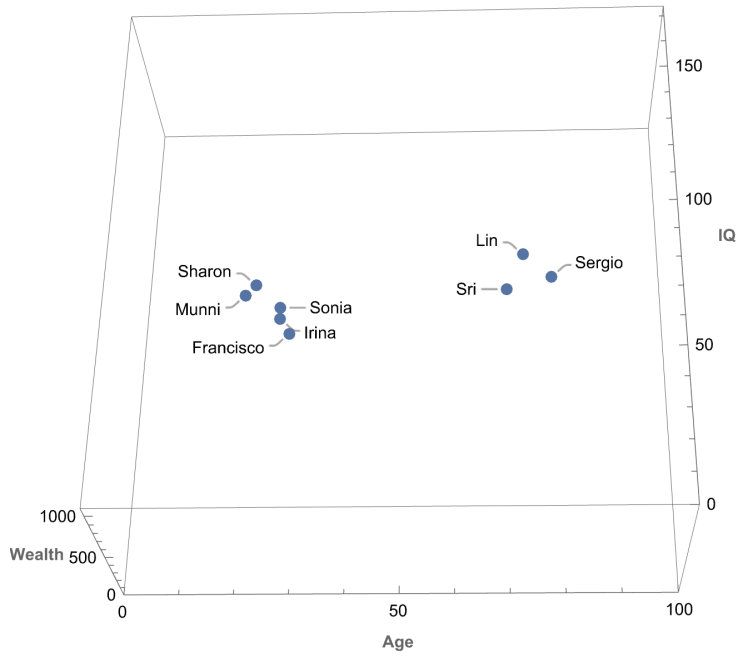

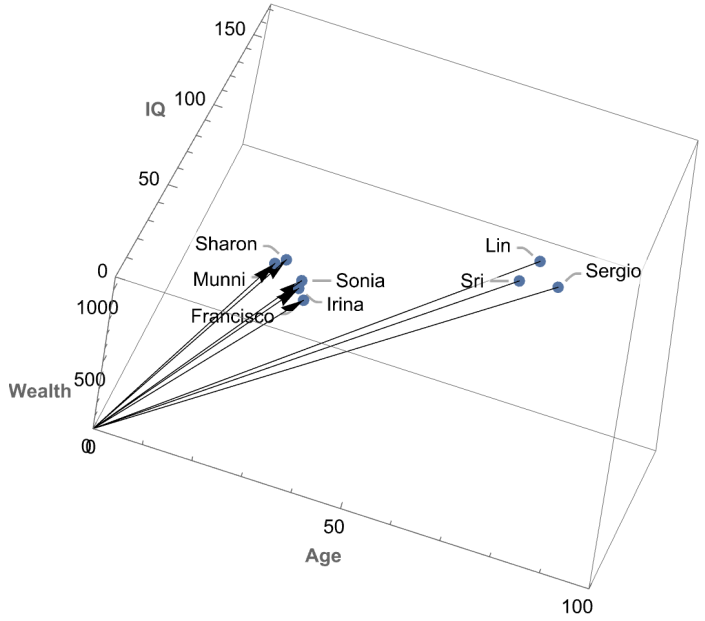

- *What story is this plot conveying?*

There are clearly two groups (technically called “clusters”).
Group 1: Old, rich people. Group 2: Young, middle-class people.
Both group have similar IQs.

- *Can you see the clusters in the data table?*

Maybe yes for this data, but in general you won’t be able to. That’s why plotting/visualization is an essential part of data science!

Name	Age	IQ	Wealth (k$)
Francisco	28	92	16
Irina	26	98	10
Sergio	78	84	800
Munni	19	107	17
Lin	73	90	860
Sharon	21	111	18
Sonia	26	102	20
Sri	70	79	822

Tabular Data

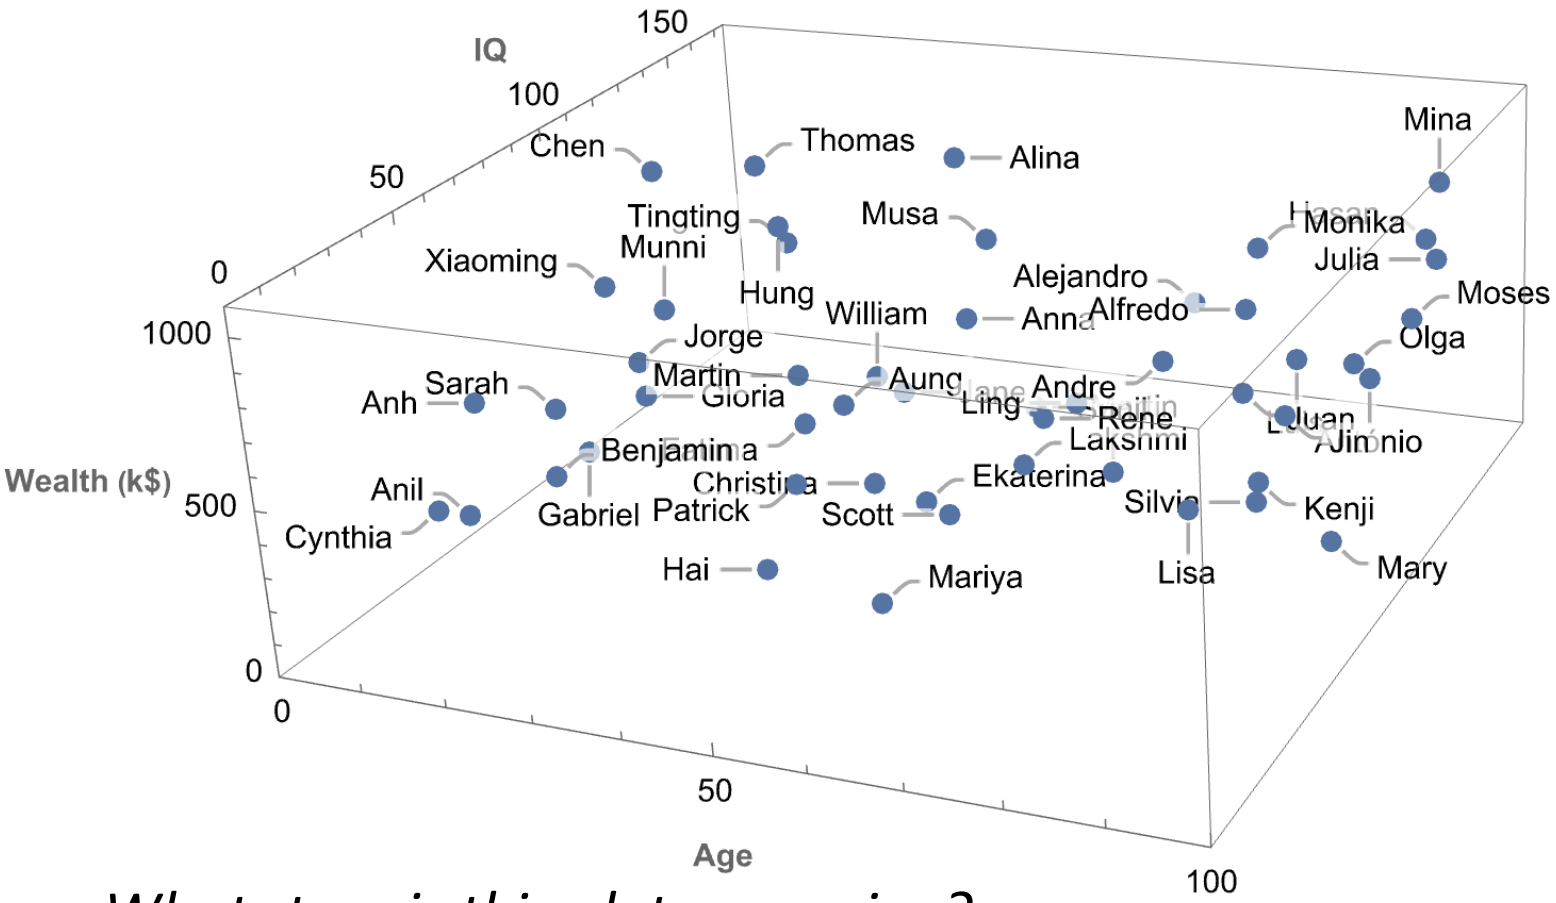

Name	Age	IQ	Wealth (k$)
Sergey	64	62	146
Dilip	52	73	952
Hai	63	153	585
Zhi	38	99	212
Patrick	55	145	522
Xing	57	76	448
Di	44	50	442
Linh	91	77	600
Lucia	12	160	152
Pedro	6	150	578
Mona	1	42	222
Margaret	96	75	481
Valentina	10	147	890
Sunil	21	119	708
Alfredo	100	123	897
Urmila	71	147	738
Tatyana	76	141	404

- *What story is this plot conveying?*

Well, this dataset looks like randomly distributed in all directions without clear clusters or correlations or trends. It looks random to eye; maybe there are structures we can't see by eye.

Tabular Data

- Data itself may be 3D (=has 3 columns) but it maybe living on a 2D (plane), 1D (line) or 0D (point) subspace . The data below is actually living on a two dimensional plane. The coordinates of this plane is a mixture of Age, IQ and Wealth. In the following weeks we'll do "dimensional reduction" to get rid of extra dimensions for easier analysis and interpretation of data.

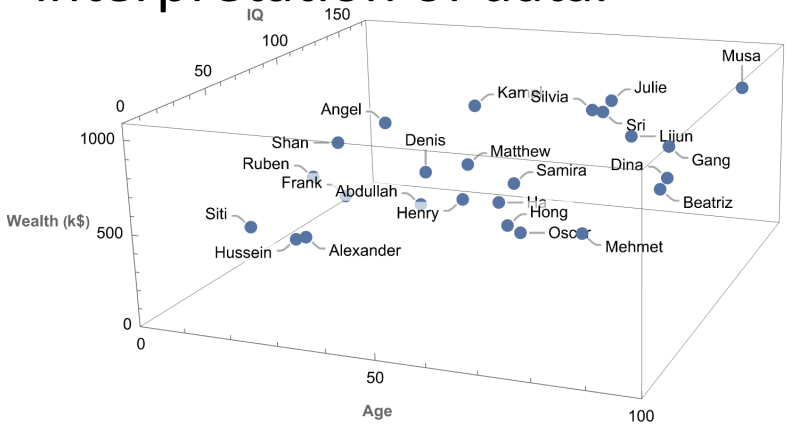


```
Plot3D[5 x + 3 y, {x, 0, 100}, {y, 0, 150}, AxesLabel → {"x", "y", "z"}]
```

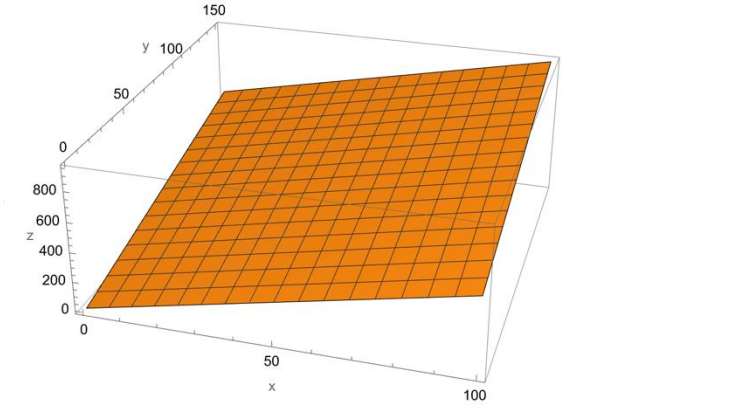

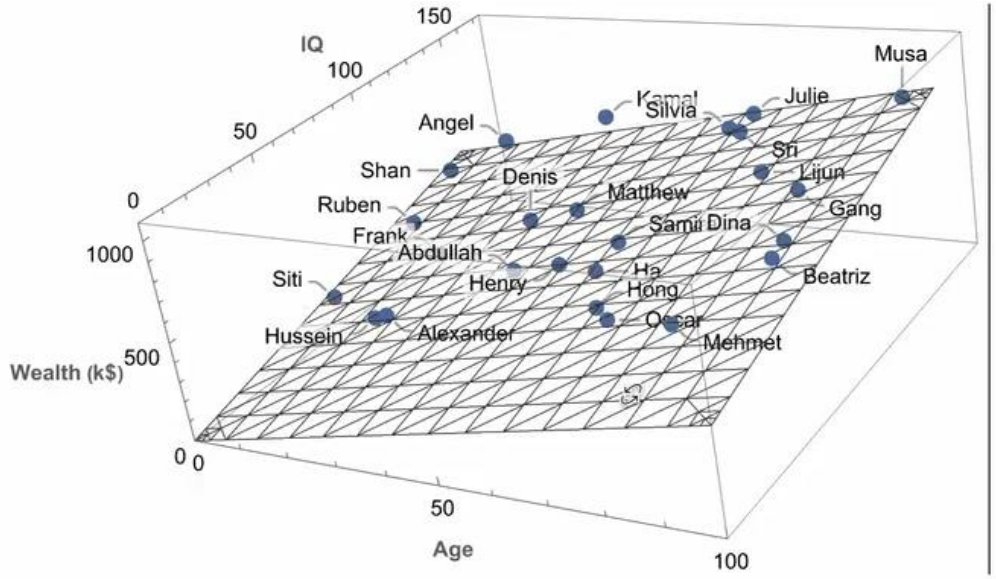

Tabular Data

- Real world data has more than 3 dimensions. Human brain cannot imagine spaces with more than 3 dimensions.

- Sample high dimensional EHR (electronic health records) dataset:

Patient ID	Full Name	Age	Gender	Blood Type	Diagnosis	Medications	Last Visit Date	Heart Rate (bpm)	Blood Pressure	BMI
1	John Doe	45	M	O+	Hypertension	Lisinopril	2023–01–15	78	150/95	27.5
2	Jane Smith	29	F	A–	Type 2 Diabetes	Metformin	2023–02–10	72	120/80	23.1
3	Robert Brown	67	M	B+	Osteoarthritis	Ibuprofen	2023–02–20	75	140/85	28.2
4	Emily Johnson	34	F	AB+	Depression	Fluoxetine	2023–03–05	70	115/75	22.4
5	Michael Williams	50	M	O–	Asthma	Albuterol Inhaler	2023–03–15	76	130/82	26.7
6	Sarah Jones	40	F	A+	Migraine	Sumatriptan	2023–04–01	74	125/78	24.9
7	William Davis	60	M	B–	Chronic Bronchitis	Azithromycin	2023–04–10	77	135/88	27.8
8	Jessica Garcia	28	F	AB–	Anemia	Iron Supplements	2023–05–05	71	110/70	21.6
9	David Martinez	72	M	A+	Rheumatoid Arthritis	Methotrexate	2023–05–15	73	145/90	28.0
10	Angela White	38	F	O+	Hypothyroidism	Levothyroxine	2023–06–01	75	128/80	25.3

Text data

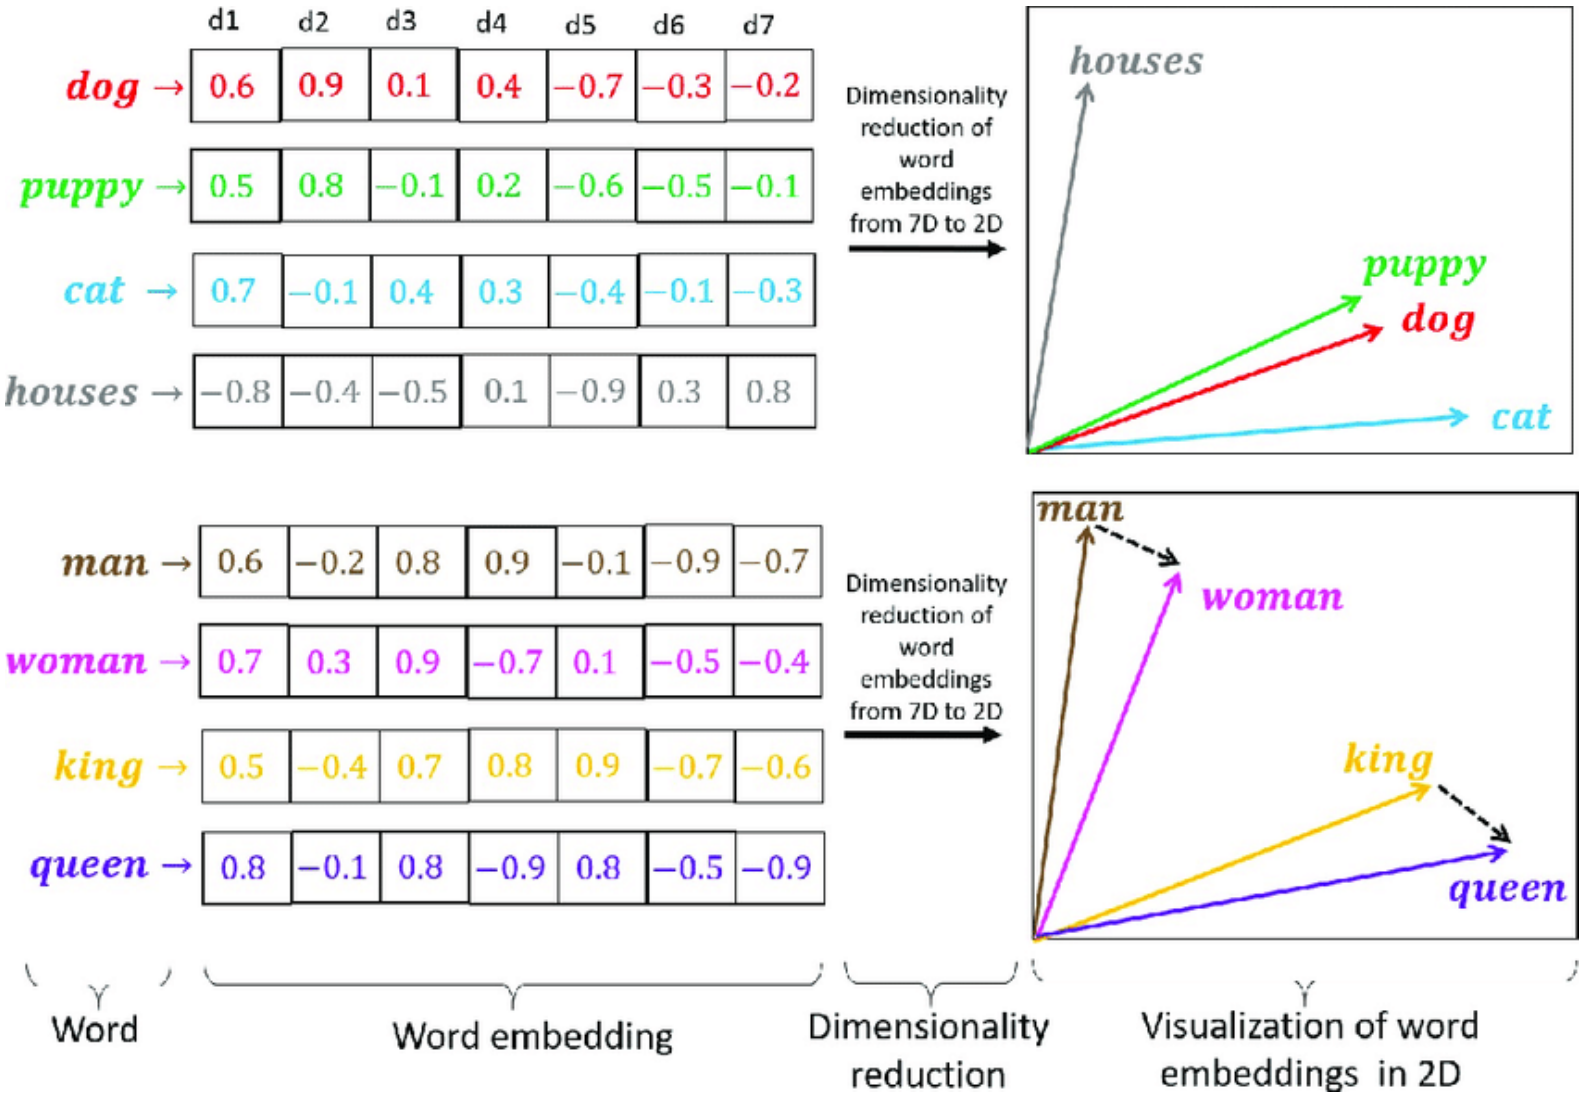
d1 d2 d3 d4 d5 d6 d7
dog → 0.6 0.9 0.1 0.4 −0.7 −0.3 −0.2
puppy → 0.5 0.8 −0.1 0.2 −0.6 −0.5 −0.1
cat → 0.7 −0.1 0.4 0.3 −0.4 −0.1 −0.3
houses → −0.8 −0.4 −0.5 0.1 −0.9 0.3 0.8
man → 0.6 −0.2 0.8 0.9 −0.1 −0.9 −0.7
woman → 0.7 0.3 0.9 −0.7 0.1 −0.5 −0.4
king → 0.5 −0.4 0.7 0.8 0.9 −0.7 −0.6
queen → 0.8 −0.1 0.8 −0.9 0.8 −0.5 −0.9
Dimensionality reduction of word embeddings from 7D to 2D
Dimensionality reduction of word embeddings from 7D to 2D
houses
puppy
dog
cat
man
woman
king
queen
Word
Word embedding
Dimensionality reduction
Visualization of word embeddings in 2D

Image data

Grayscale images are matrices (2D).

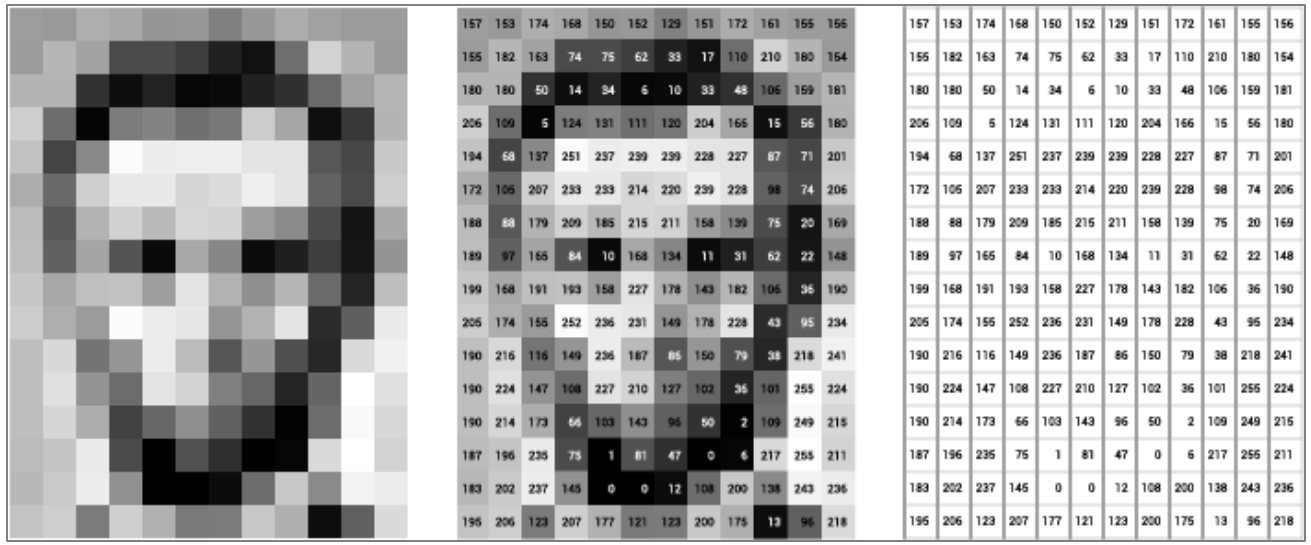

RGB colored images are rank-3 tensors (3D).

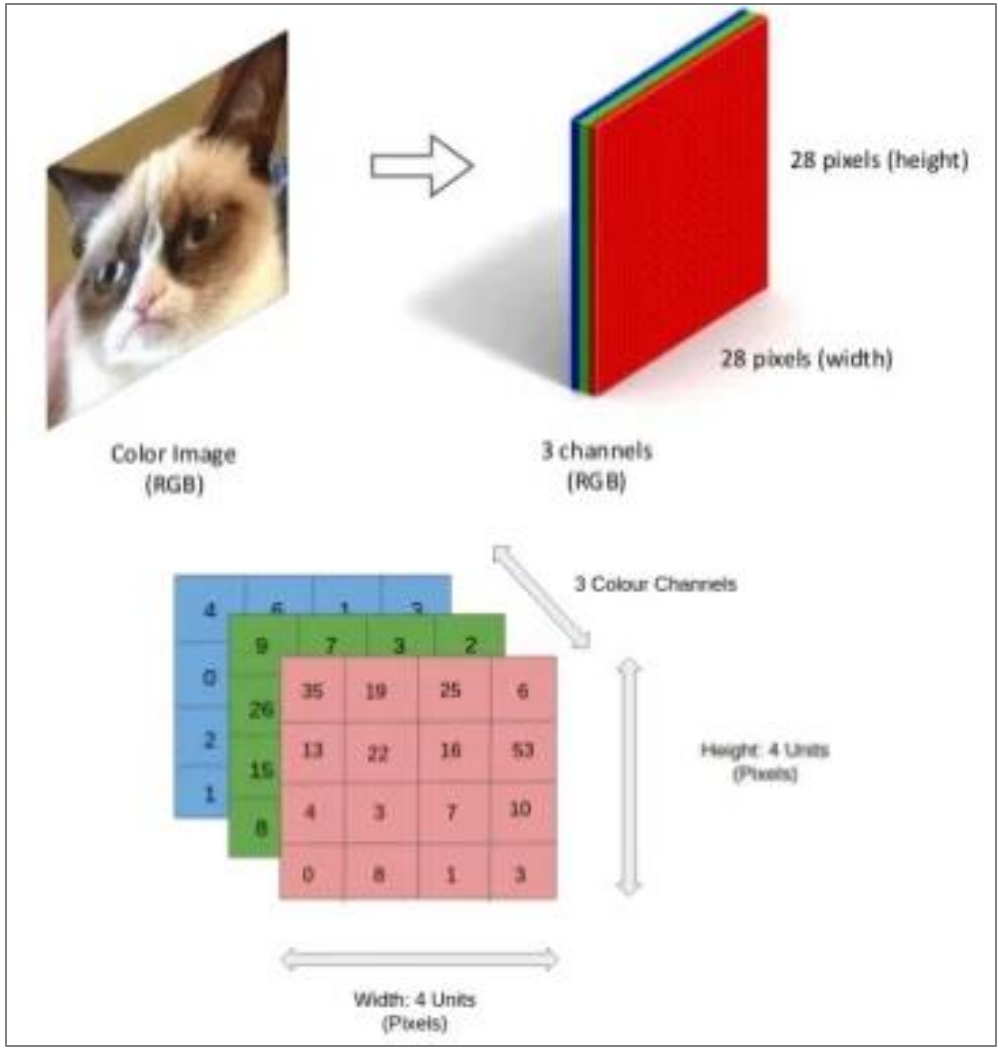

Video data

- Video data is a rank-4 tensor (4D) with dimensions:

number of frames (time) x height x width x color channels.